

RSA® Conference 2019

San Francisco | March 4–8 | Moscone Center



BETTER.

SESSION ID: ASD-W10

Practical Approaches to Cloud Native Security

Karthik Gaekwad

Principal Engineer
Oracle Inc
@iteration1



#RSAC

Hello

- I'm Karthik Gaekwad
- NOT a DBA



- <https://cloudnative.oracle.com/>
- Cloud Native Evangelist at Oracle Cloud
- Past: Developer on the Oracle Managed Kubernetes Team

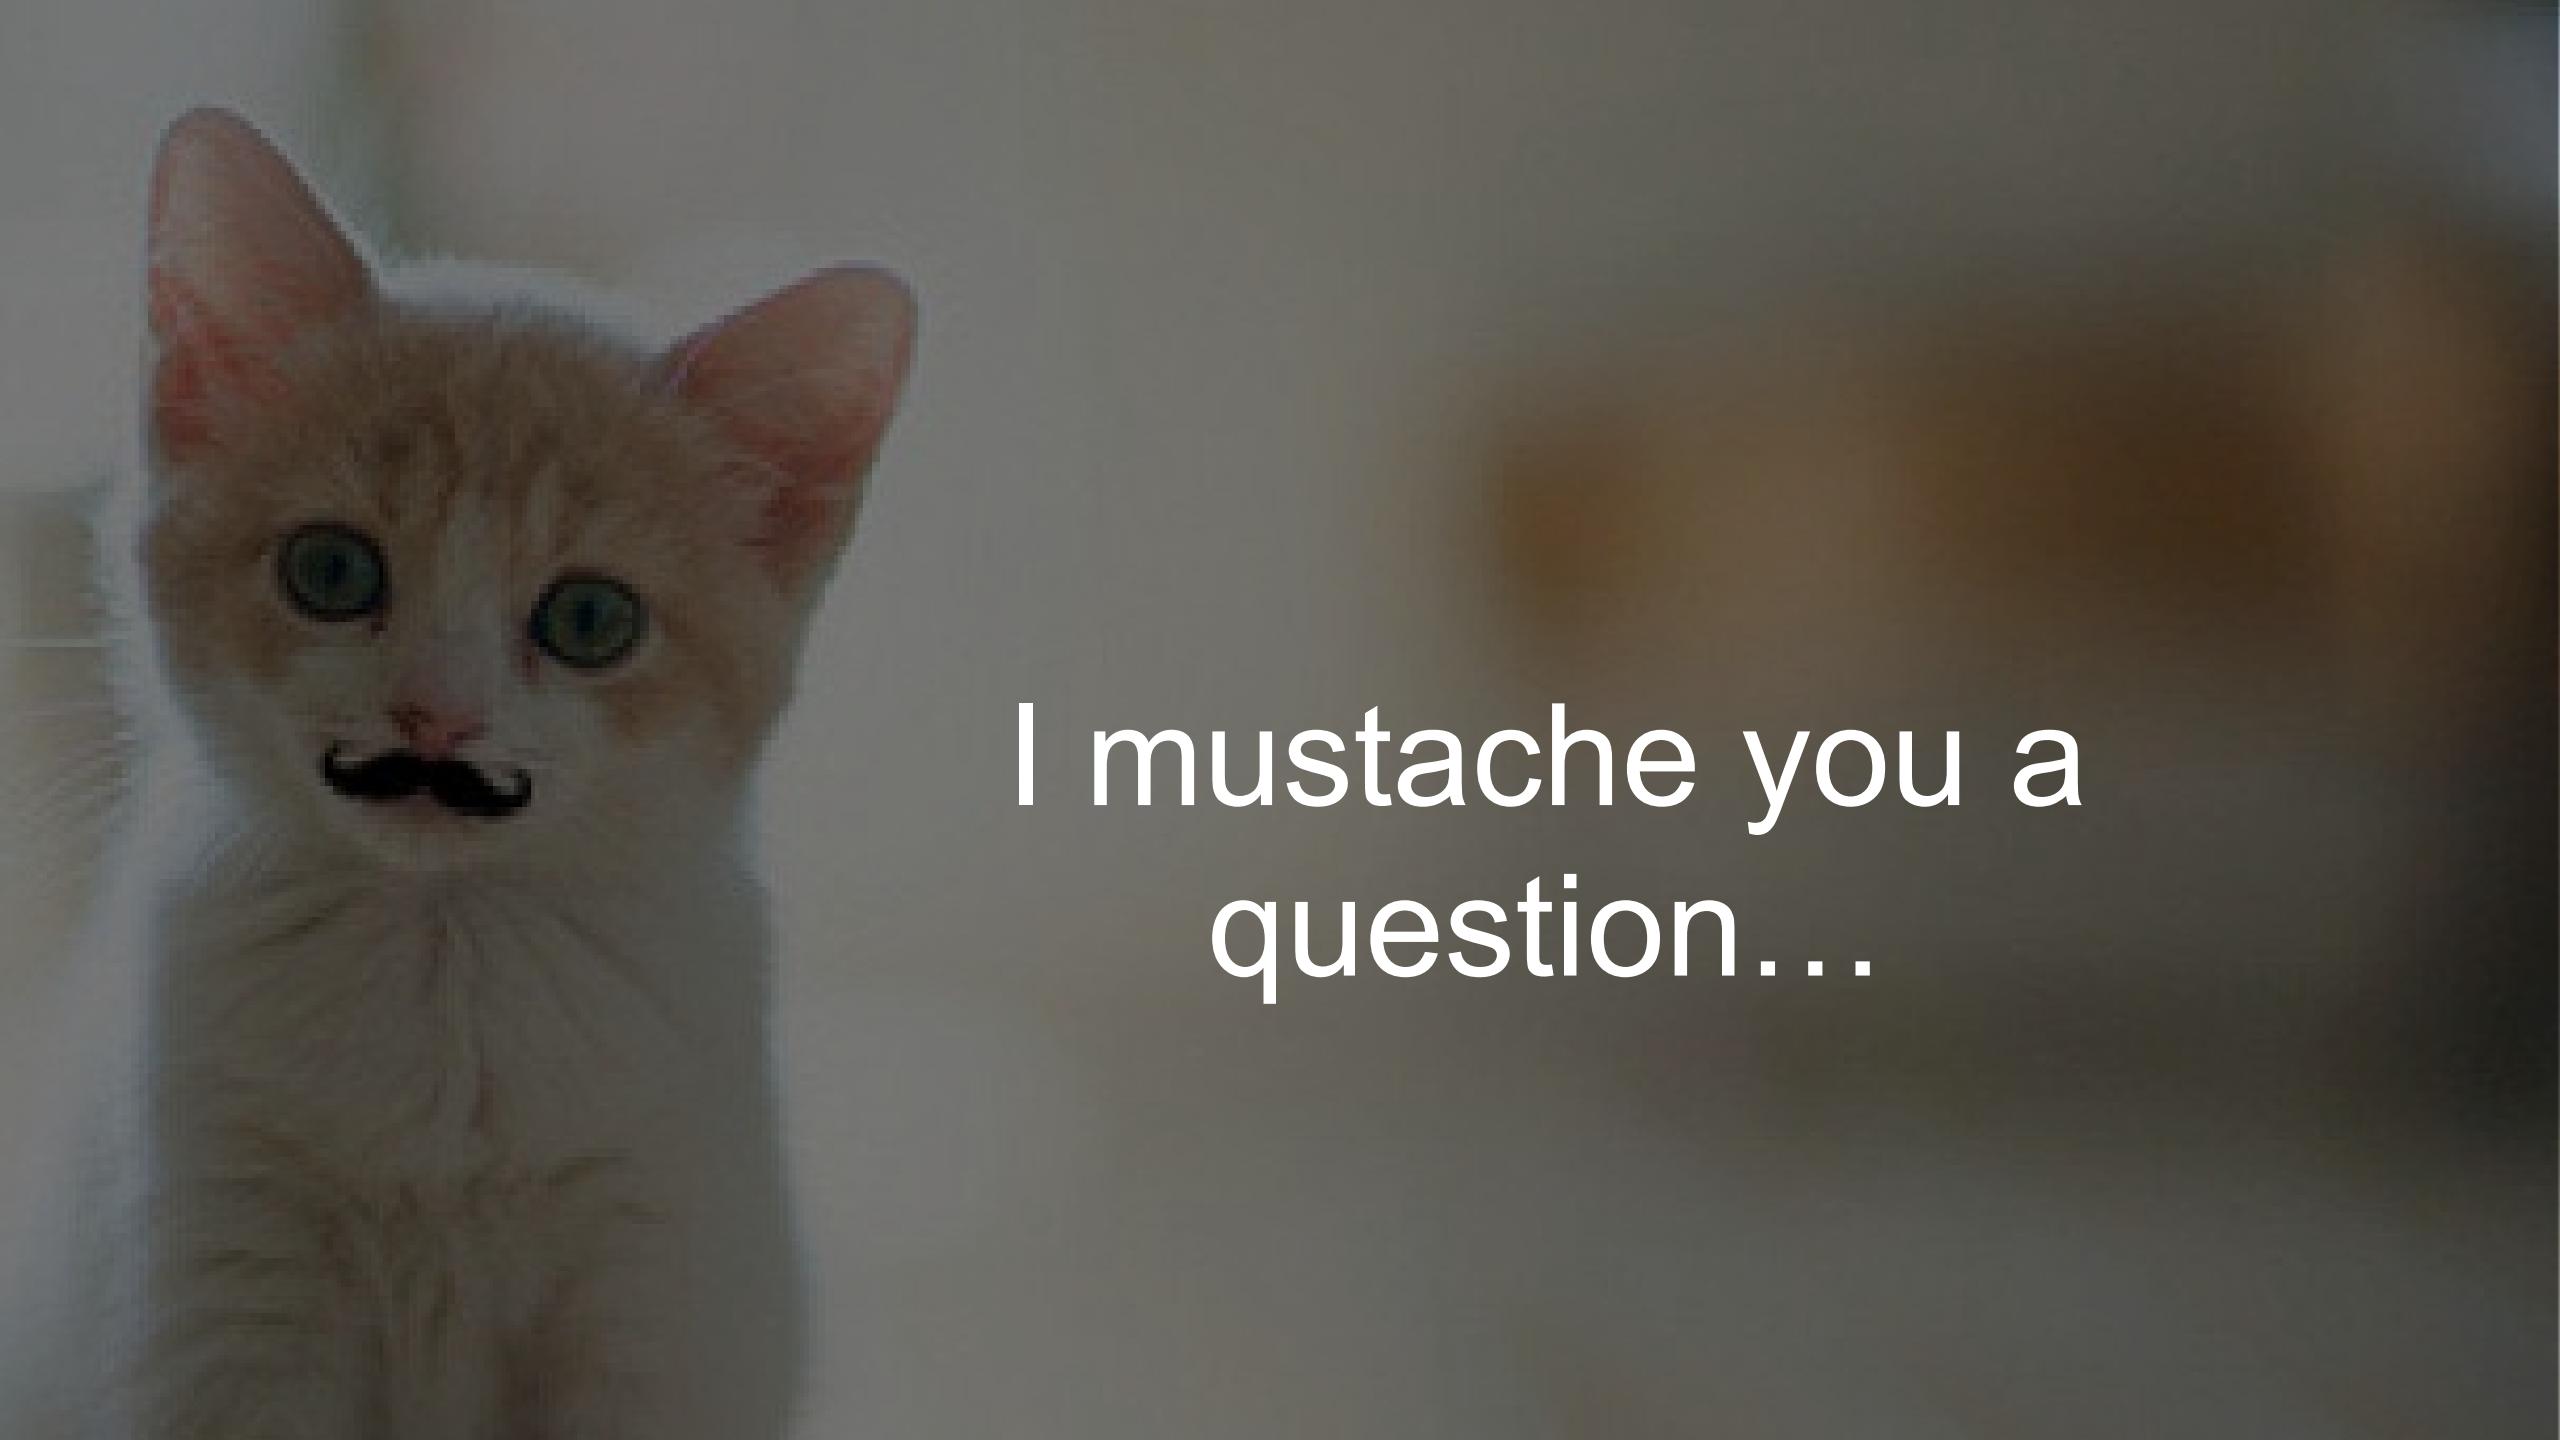


@iteration1

Hello



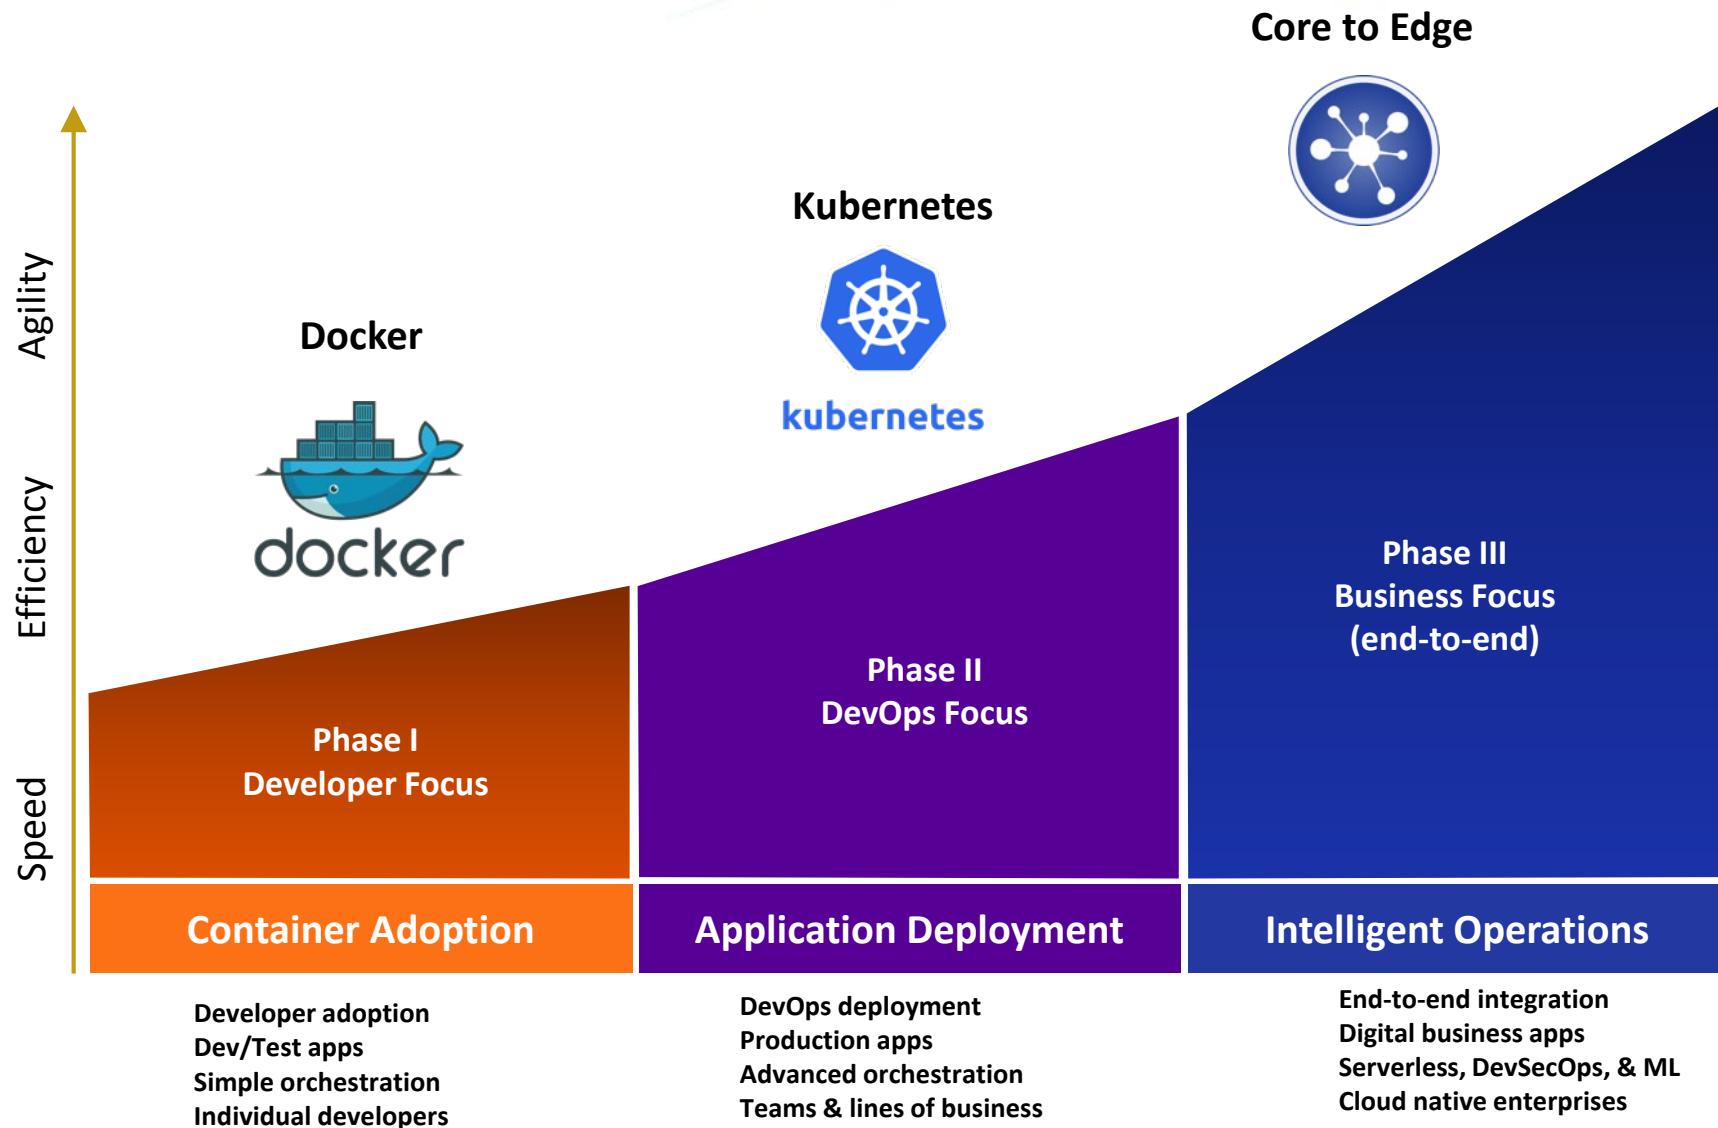
- Been in Industry 15 years.
- In general, I like building stuff with friends.
 - Maintainer for Gauntlet- Open source security scanner.
- Love Teaching and building community.
 - Run Devopsdays Austin, Container Days, Cloud Austin.
 - Chair All Day Devops Cloud Native track.
 - LinkedIn Learning Author for Learning Kubernetes (and more).

A close-up photograph of a young, light orange or cream-colored kitten. The kitten has large, expressive green eyes and a small, dark, curved mustache on its pink nose. It is looking directly at the camera with a curious expression. The background is a soft, out-of-focus grey.

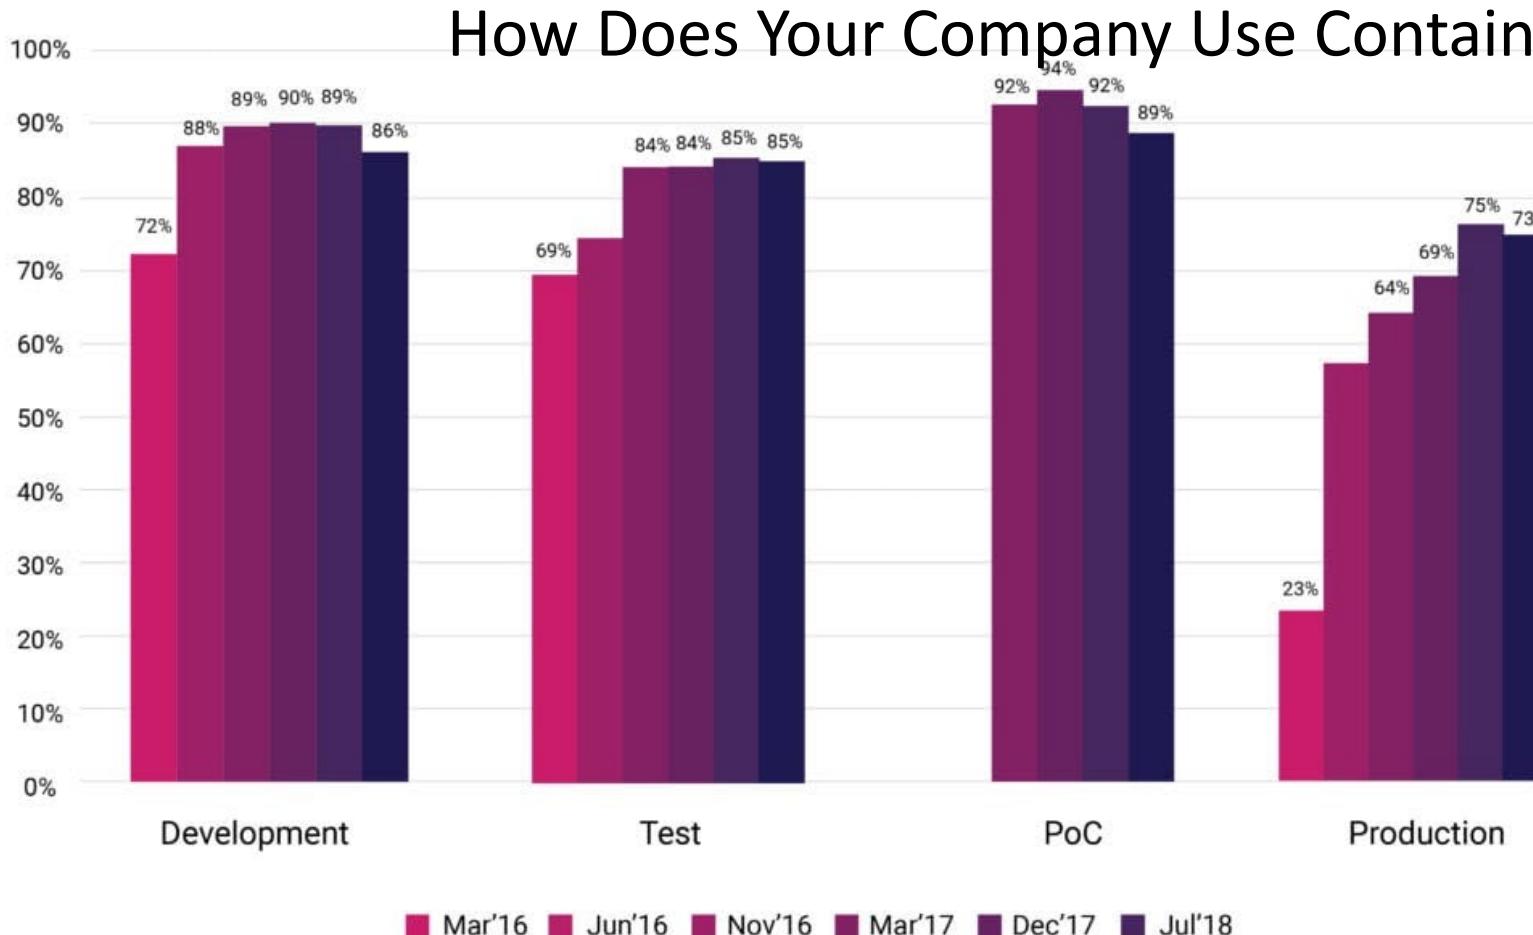
I mustache you a
question...

The Cloud Native Journey

#RSAC



CNCF Survey: August 2018

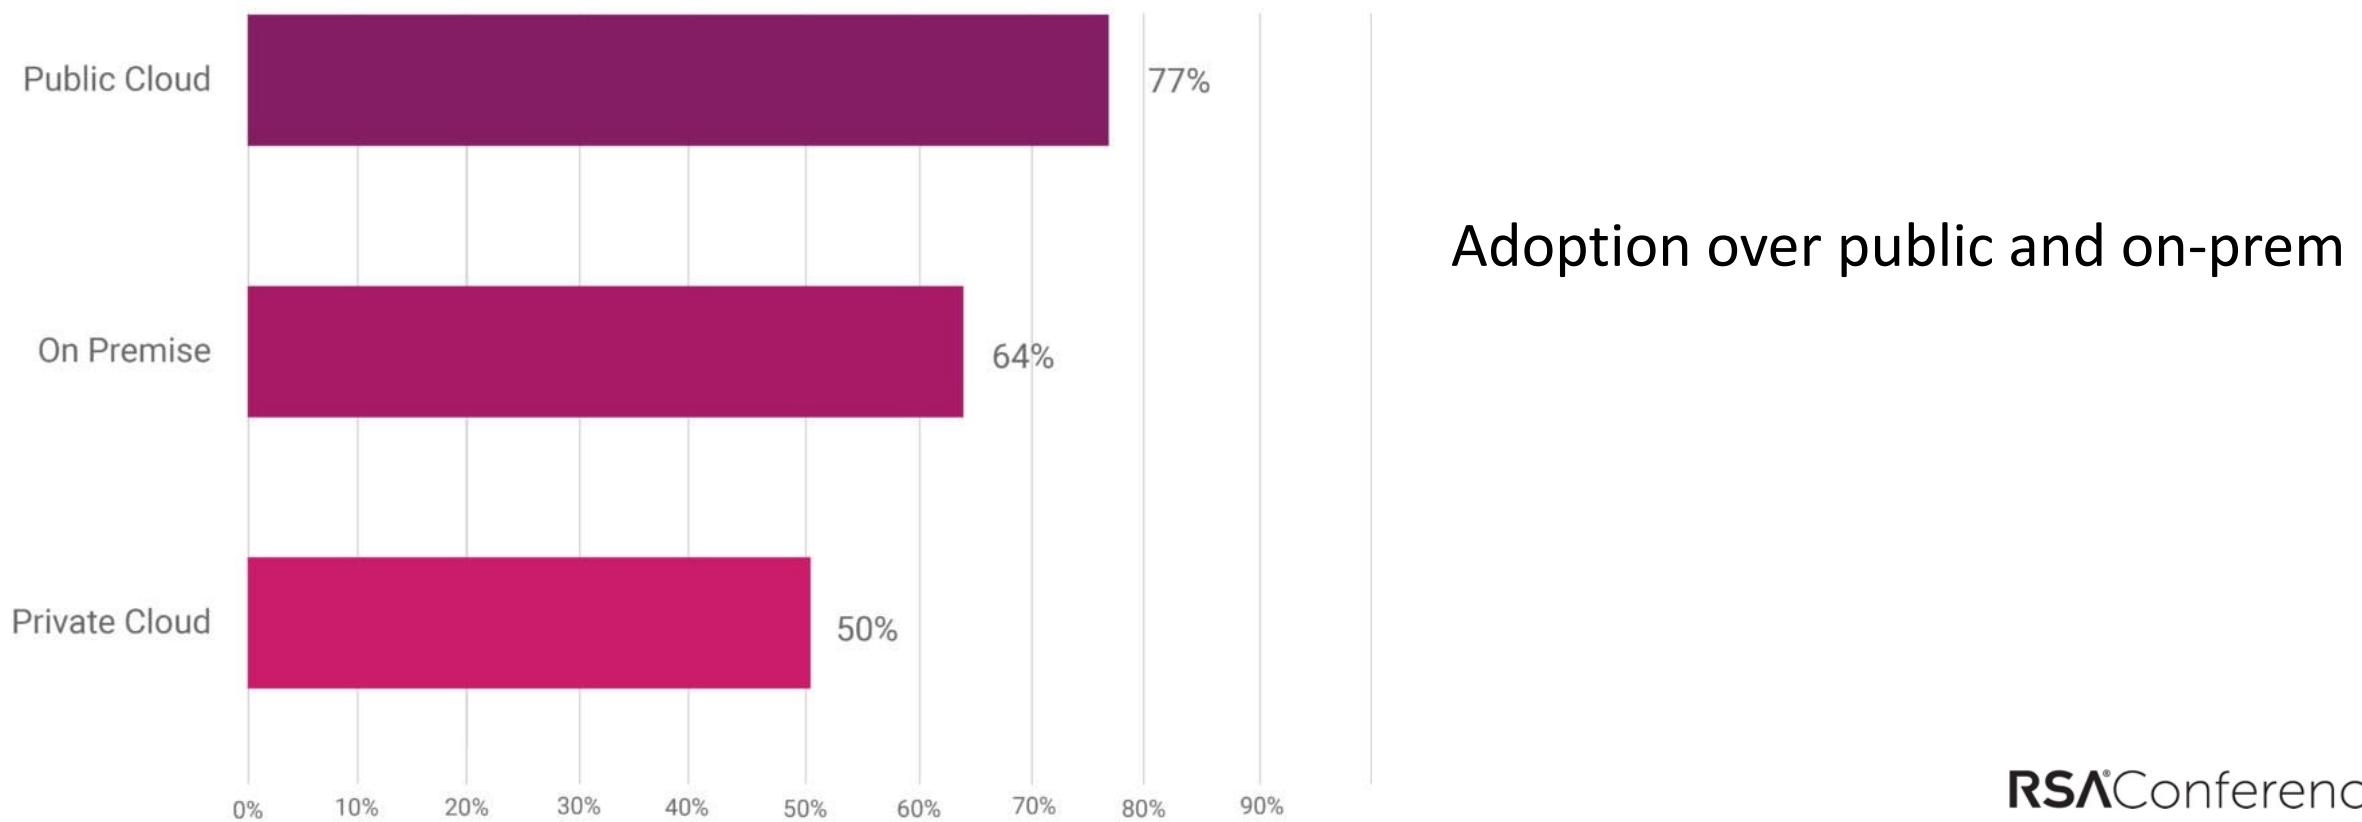


Lots of adoption on dev/staging

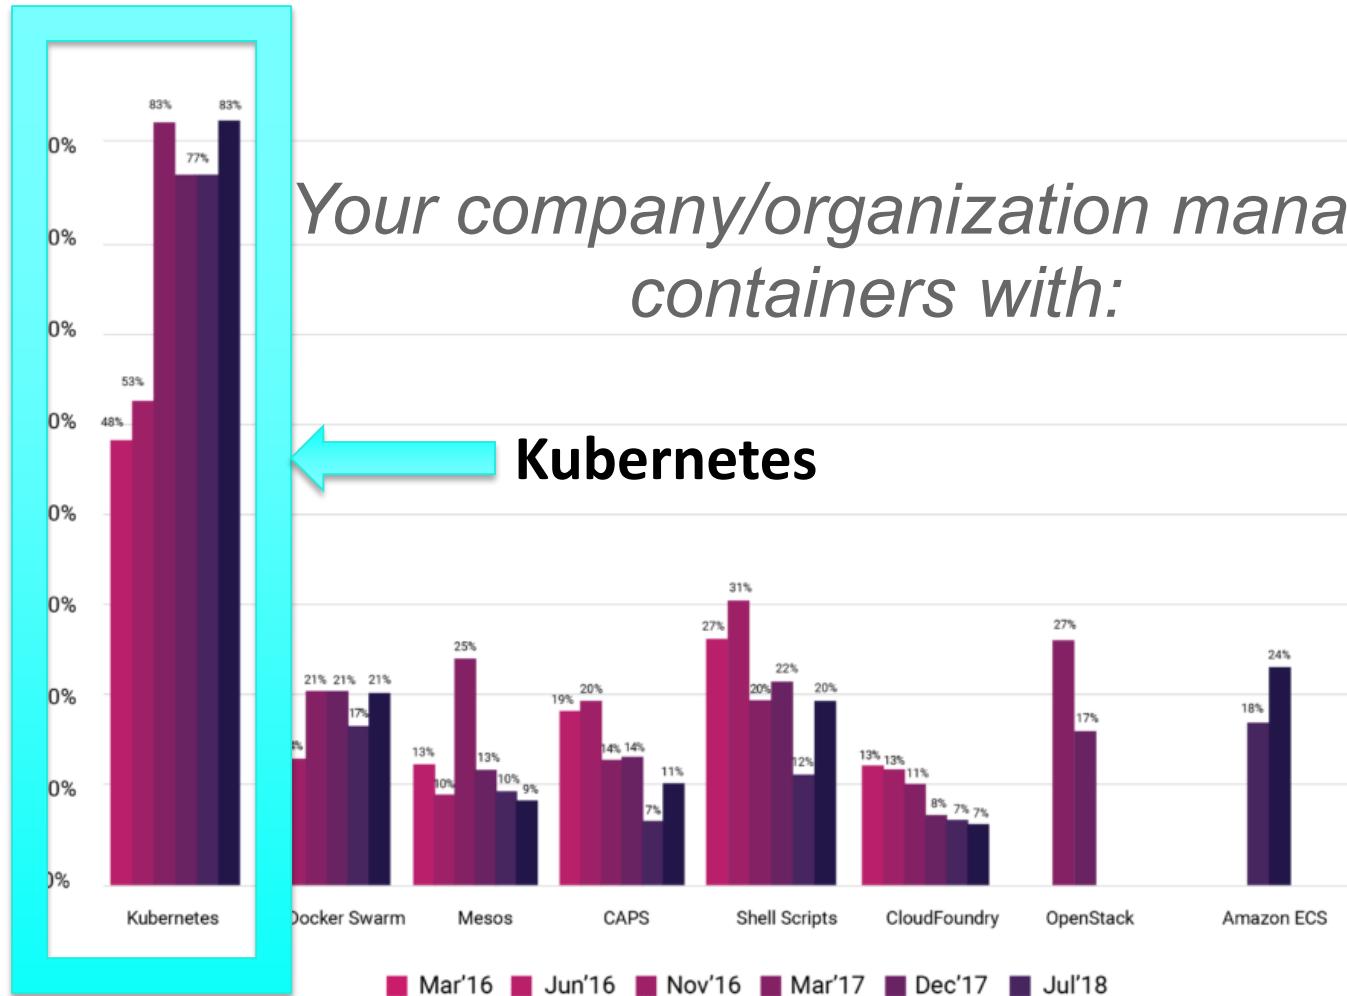
Continued production increase

CNCF Survey: August 2018

How Does Your Company Use Containers and Where?



Kubernetes Dominates Container Management



Your company/organization manages containers with:

Kubernetes

Good News, Bad News...

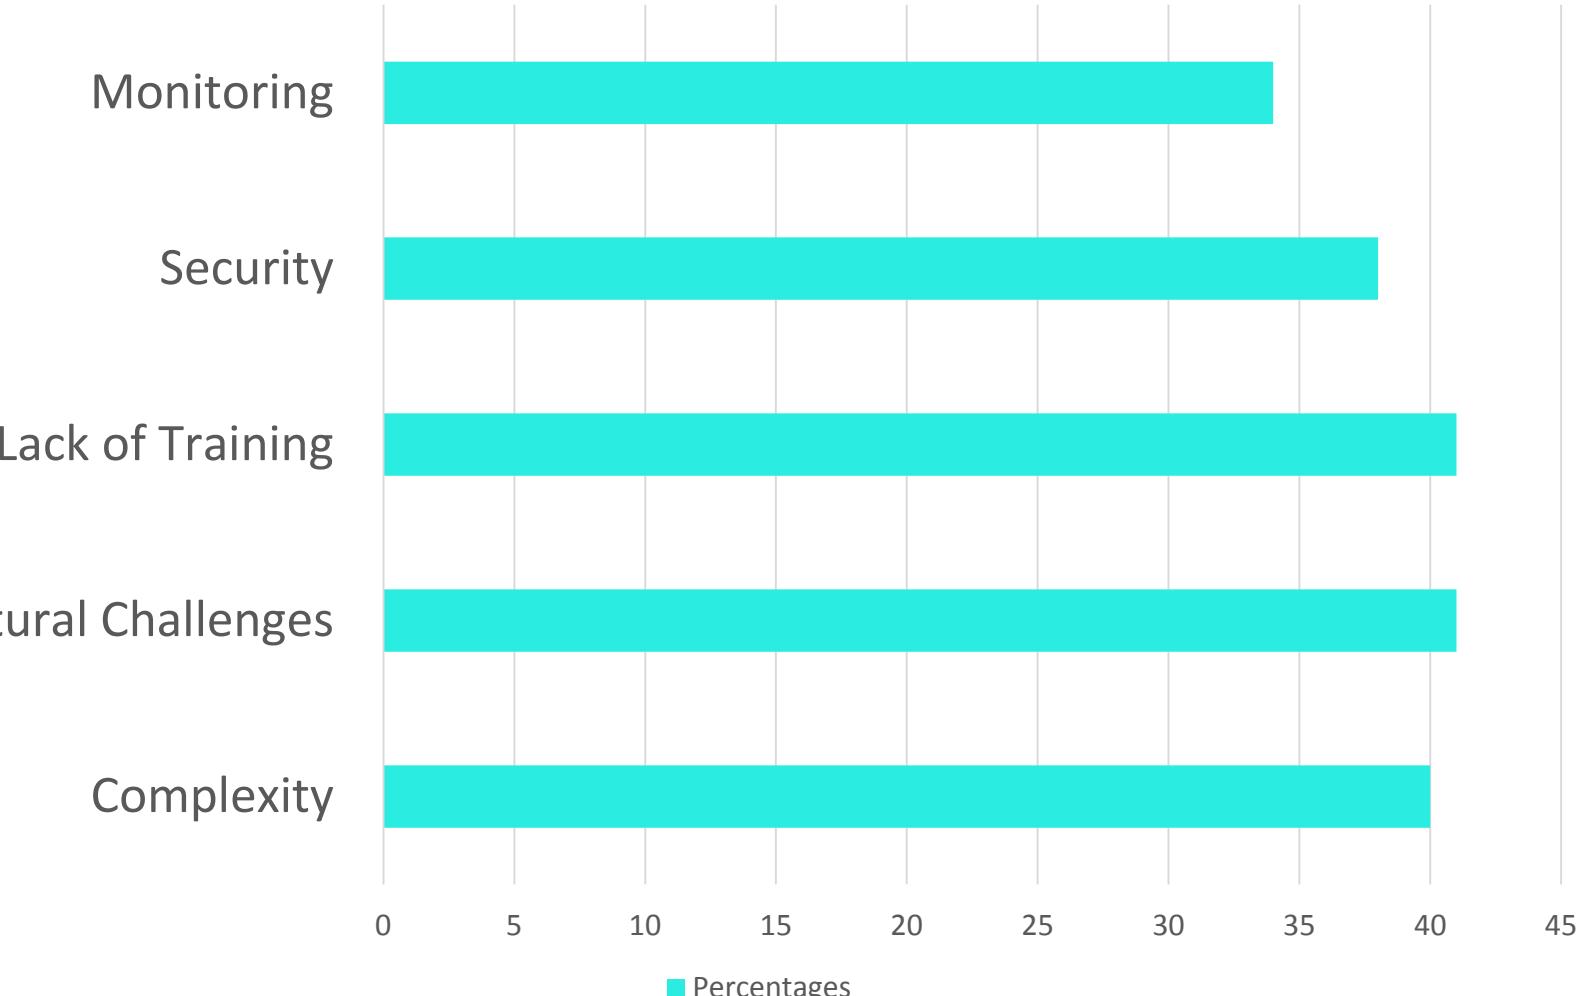
Many Projects...



Good usage in dev/prod

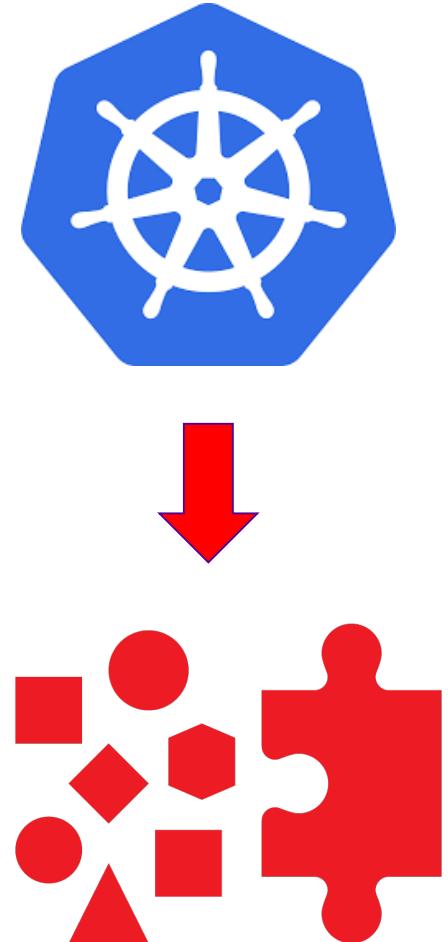
But...

Top 5 challenges to cloud native adoption...



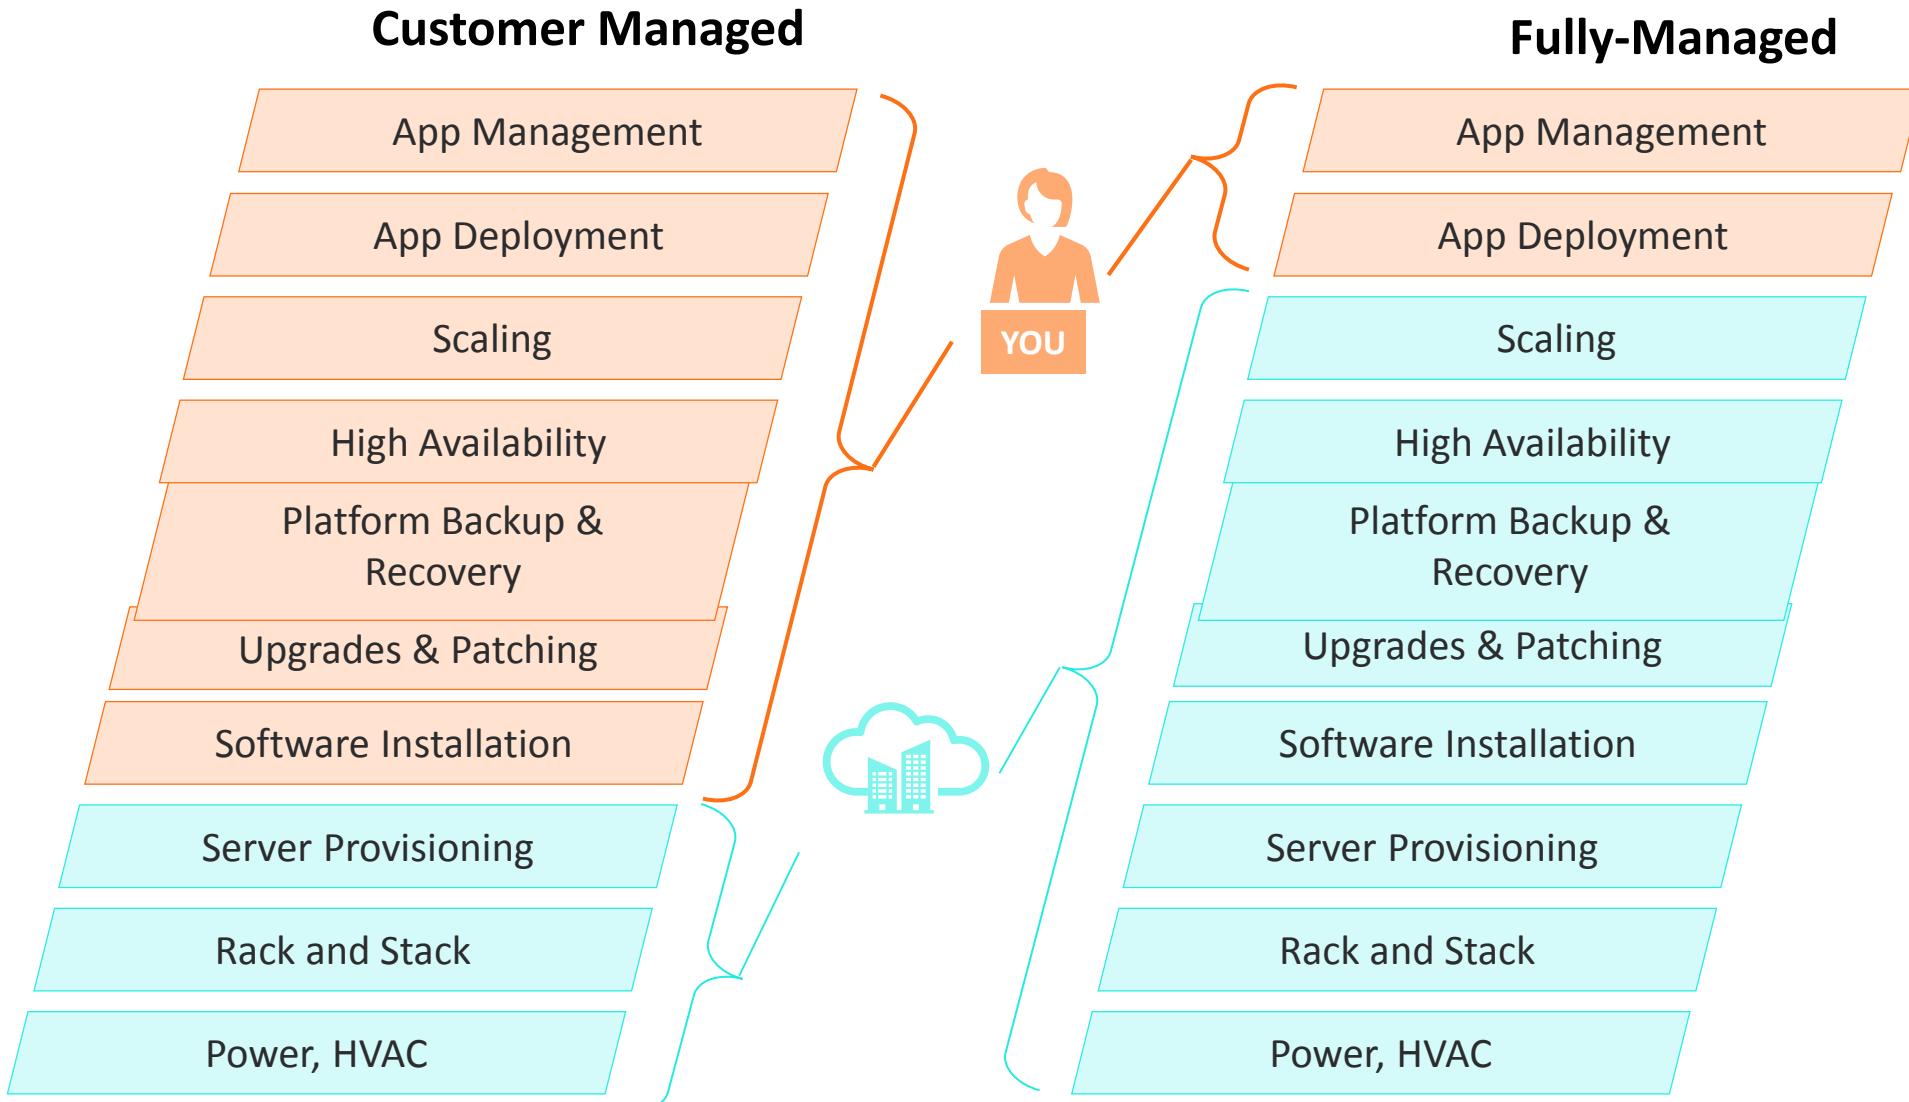
Kubernetes & Cloud Native Challenges

- Managing, maintaining, upgrading Kubernetes Control Plane
 - API Server, etcd, scheduler etc....
- Managing, maintaining, upgrading Kubernetes Data Plane
 - In place upgrades, deploy parallel cluster etc....
- Figuring out container networking & storage
 - Overlays, persistent storage etc... - it should just work
- Managing Teams
 - How do I manage & control team access to my clusters?
- Security, security, security



Source: Oracle Customer Survey 2018

How Are Teams Addressing Complexity, Training Issues?



Benefits

- ✓ Faster Time to Deploy
- ✓ Lower Risk
- ✓ Accelerate Innovation

A close-up shot of a man with light brown, curly hair and blue eyes. He is wearing a dark brown top hat and a purple velvet jacket over a green shirt. He has a wide, joyful smile and is resting his chin on his hand, looking slightly off-camera to the right.

Which brings us to security...



DAYS WITHOUT
AN ACCIDENT

Where no news, is good news!

Unsecured K8s dashboards



Lessons from the Cryptojacking Attack at Tesla

by RedLock CSI Team | 02.20.18, 6:00 AM

- Unsecured Kubernetes Dashboard with account creds.
- Used this to mine cryptocurrency.
- 2017: Aviva
- 2018: Tesla, Weight Watchers

- <https://redlock.io/blog/cryptojacking-tesla>

Kubelet credentials hack

The screenshot shows a HackerOne report for issue #341876. The report details a Server-Side Request Forgery (SSRF) vulnerability in Shopify's Exchange service, which led to root access in all instances. The report was resolved on May 23, 2018, at 4:09pm -0500. The bounty was \$25,000. The exploit chain involved creating a store on partners.shopify.com and editing the password.liquid template to redirect to Google Cloud Metadata. The exploit code is shown in a code block:

```

<script>
window.location="http://metadata.google.internal/computeMetadata/v1beta1/instance/service-accounts/default
// iframes don't work here because Google Cloud sets the 'X-Frame-Options: SAMEORIGIN' header.
</script>

```

- Shopify: Server Side request Forgery
- Get kubelet certs/private key
- Root access to any container in part of infrastructure.
- <https://hackerone.com/reports/341876>

← → C https://www.shodan.io/search?query=port%3A"2379"+product%3A"etcd"

SHODAN
port:"2379" product:"etcd"

Explore
Downloads
Reports
Developer Pricing
Enterprise Access

Exploits
Maps
Share Search
[Download Results](#)
[Create Report](#)

TOTAL RESULTS

2,367

TOP COUNTRIES



China	933
United States	602
Germany	153
France	110
Singapore	67

TOP ORGANIZATIONS

Amazon.com	323
Hangzhou Alibaba Advertising Co.,Ltd.	276
Tencent cloud computing	207
Hetzner Online GmbH	70
Digital Ocean	70

TOP OPERATING SYSTEMS

Linux 3.x	1
-----------	---

TOP VERSIONS

3.3.2	299
3.2.18	138
3.3.9	128
3.2.22	123
3.2.17	82

110.24.128.101

Tencent cloud computing
Added on 2018-09-26 19:15:25 GMT

China
[Details](#)

etcd
Name: etcd_10.0.128.13
Version: 3.3.2
Uptime: 79h30m33.232103055s
Peers: http://10.0.128.13:2380

10.236.207.66

2.compute.amazonaws.com
Amazon Corporate Services Pty
Added on 2018-09-26 19:08:40 GMT
 Australia, Sydney

[Details](#)

etcd
Name: NFR-50nodeap-southeast-2002
Version: 3.3.2
Uptime: 54h22m1.357953836s
Peers: http://13.236.207.66:2380

50.84.104.94

Spectrum Business
Added on 2018-09-26 18:52:54 GMT
 United States, Flower Mound
[Details](#)

etcd
Name: core2
Version: 2.2.5
Uptime: 5h38m59.402175596s
Peers: http://10.11.36.2:2380, http://10.11.36.2:7001

192.168.9.100

Red Hat
Added on 2018-09-26 18:49:53 GMT
 United States
[Details](#)

etcd
Name: rdocloud-devstack4.rdocloud
Version: 3.2.17
Uptime: 633h2m38.382059797s
Peers: http://192.168.4.6:2380

54.195.264.221

Amazon.com
Added on 2018-09-26 18:45:20 GMT
 United States, Ashburn
[Details](#)

etcd
Name: master-0
Version: 3.2.18
Uptime: 20m19.490052863s

SHODAN port:"2379" product:"etcd"   Explore Downloads Reports Developer Pricing Enterprise Access

TOTAL RESULTS 2,367



Country	Count
China	933
United States	602
Germany	153
France	110
Singapore	67

TOP ORGANIZATIONS

Organization	Count
Amazon.com	323
Hangzhou Alibaba Advertising Co.,Ltd.	276
Tencent cloud computing	207
Hetzner Online GmbH	70
Digital Ocean	70

TOP OPERATING SYSTEMS

Operating System	Count
Linux 3.x	1

TOP VERSIONS

Version	Count
3.3.2	299
3.2.18	138
3.3.9	128
3.2.22	123
3.2.17	82

110.24.128.101
Tencent cloud computing
Added on 2018-09-26 19:15:25 GMT
 China
[Details](#)

etcd
Name: etcd_10.0.128.13
Version: 3.3.2
Uptime: 79h30m33.232103055s
Peers: http://10.0.128.13:2380

13.236.207.66
2.compute.amazonaws.com
Amazon Corporate Services Pty
Added on 2018-09-26 19:08:40 GMT
 Australia, Sydney
[Details](#)

cloud

etcd
Name: NFR-50nodeap-southeast-2002
Version: 3.3.2
Uptime: 54h22m1.357953836s
Peers: http://13.236.207.66:2380

50.84.104.94
Spectrum Business
Added on 2018-09-26 18:52:54 GMT
 United States, Flower Mound
[Details](#)

etcd
Name: core2
Version: 2.2.5
Uptime: 5h38m59.402175596s
Peers: http://10.11.36.2:2380, http://10.11.36.2:7001

192.168.9.66
Red Hat
Added on 2018-09-26 18:49:53 GMT
 United States
[Details](#)

etcd
Name: rdocloud-devstack4.rdocloud
Version: 3.2.17
Uptime: 633h2m38.382059797s
Peers: http://192.168.4.6:2380

54.195.264.221
Amazon.com
Added on 2018-09-26 18:45:20 GMT
 United States, Ashburn
[Details](#)

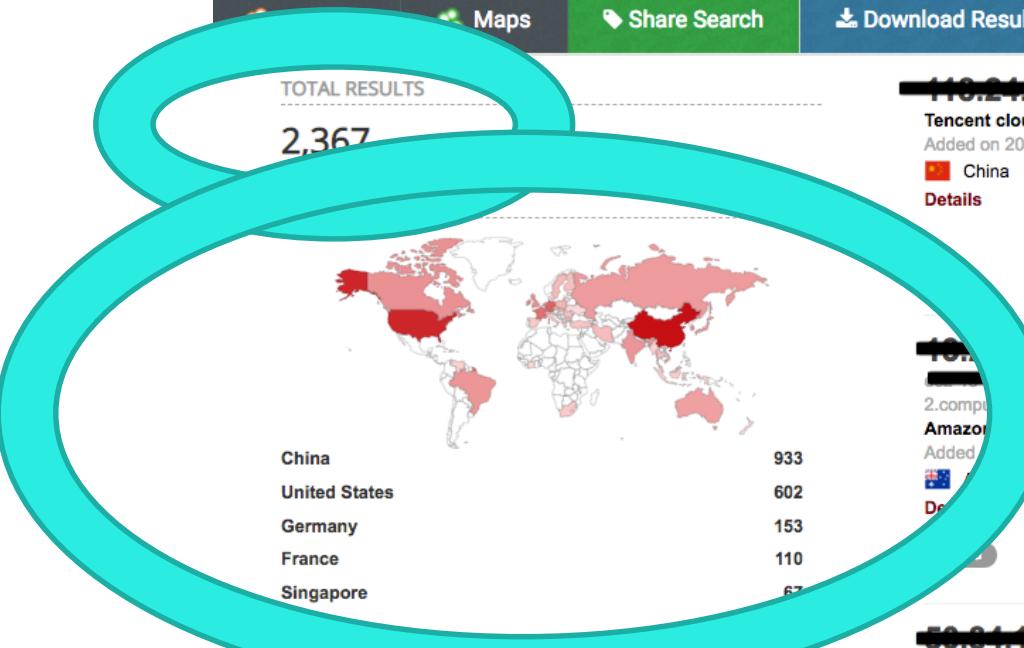
etcd
Name: master-0
Version: 3.2.18
Uptime: 20m19.490052863s

SHODAN

port:"2379" product:"etcd"

Maps Share Search Download Results Create Report

TOTAL RESULTS 2,367



China 933
United States 602
Germany 153
France 110
Singapore 67

TOP OPERATING SYSTEMS

Linux 3.x	1
-----------	---

TOP VERSIONS

3.3.2	299
3.2.18	138
3.3.9	128
3.2.22	123
3.2.17	82

110.24.128.13

Tencent cloud computing
Added on 2018-09-26 19:15:25 GMT
China
Details

13.236.207.66

2.complexity...amazonaws.com
Amazon Web Services Corporate Services Pty
Added on 2018-09-26 19:08:40 GMT
Australia, Sydney
Details

50.84.104.94

Spectrum Business
Added on 2018-09-26 18:52:54 GMT
United States, Flower Mound
Details

192.168.9.100

Red Hat
Added on 2018-09-26 18:49:53 GMT
United States
Details

54.195.264.221

Amazon.com
Added on 2018-09-26 18:45:20 GMT
United States, Ashburn
Details

etcd
Name: etcd_10.0.128.13
Version: 3.3.2
Uptime: 79h30m33.232103055s
Peers: http://10.0.128.13:2380

etcd
Name: NFR-50nodeap-southeast-2002
Version: 3.3.2
Uptime: 54h22m1.357953836s
Peers: http://13.236.207.66:2380

etcd
Name: core2
Version: 2.2.5
Uptime: 5h38m59.402175596s
Peers: http://10.11.36.2:2380, http://10.11.36.2:7001

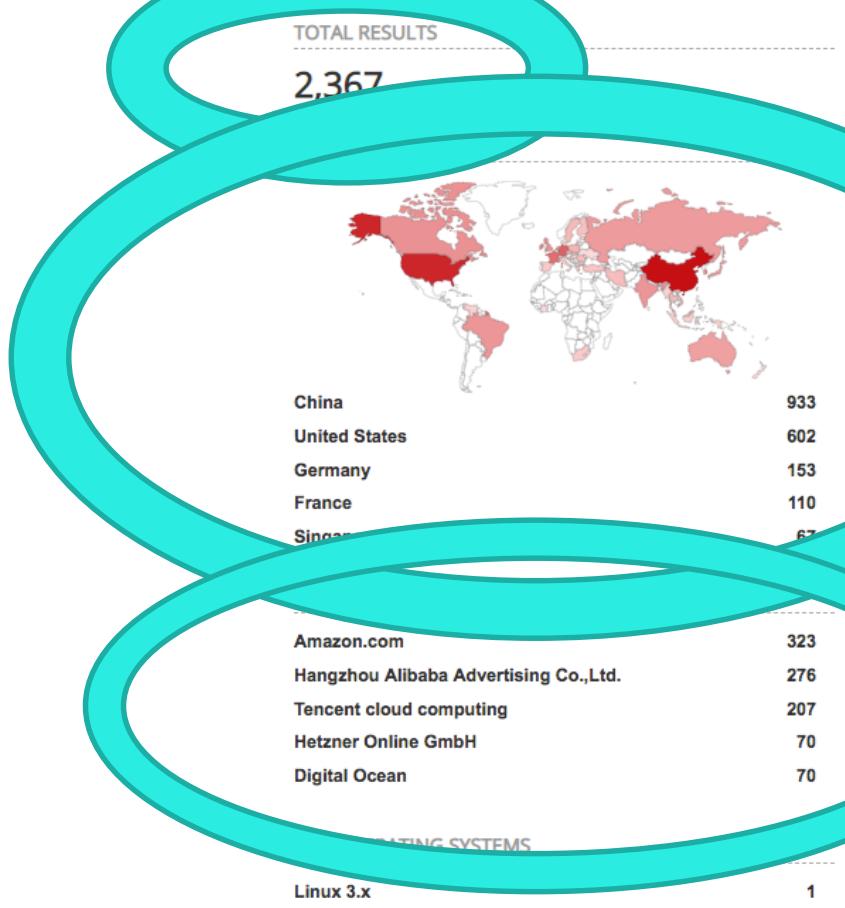
etcd
Name: rdocloud-devstack4.rdocloud
Version: 3.2.17
Uptime: 633h2m38.382059797s
Peers: http://192.168.4.6:2380

etcd
Name: master-0
Version: 3.2.18
Uptime: 20m19.490052863s

SHODAN

port:"2379" product:"etcd"

Maps Share Search Download Results Create Report



TOP VERSIONS

3.3.2	299
3.2.18	138
3.3.9	128
3.2.22	123
3.2.17	82

110.24.128.101

Tencent cloud computing
Added on 2018-09-26 19:15:25 GMT

China
[Details](#)

etcd
Name: etcd_10.0.128.13
Version: 3.3.2
Uptime: 79h30m33.232103055s
Peers: http://10.0.128.13:2380

10.236.207.66

2.com.amazonaws.com
Amazon Corporate Services Pty
Added on 2018-09-26 19:08:40 GMT
Australia, Sydney
Digital Ocean

etcd
Name: NFR-50nodeap-southeast-2002
Version: 3.3.2
Uptime: 54h22m1.357953836s
Peers: http://13.236.207.66:2380

50.84.104.94

Spectrum Business
Added on 2018-09-26 18:52:54 GMT
United States, Flower Mound
Digital Ocean

etcd
Name: core2
Version: 2.2.5
Uptime: 5h38m59.402175596s
Peers: http://10.11.36.2:2380, http://10.11.36.2:7001

192.168.9.100

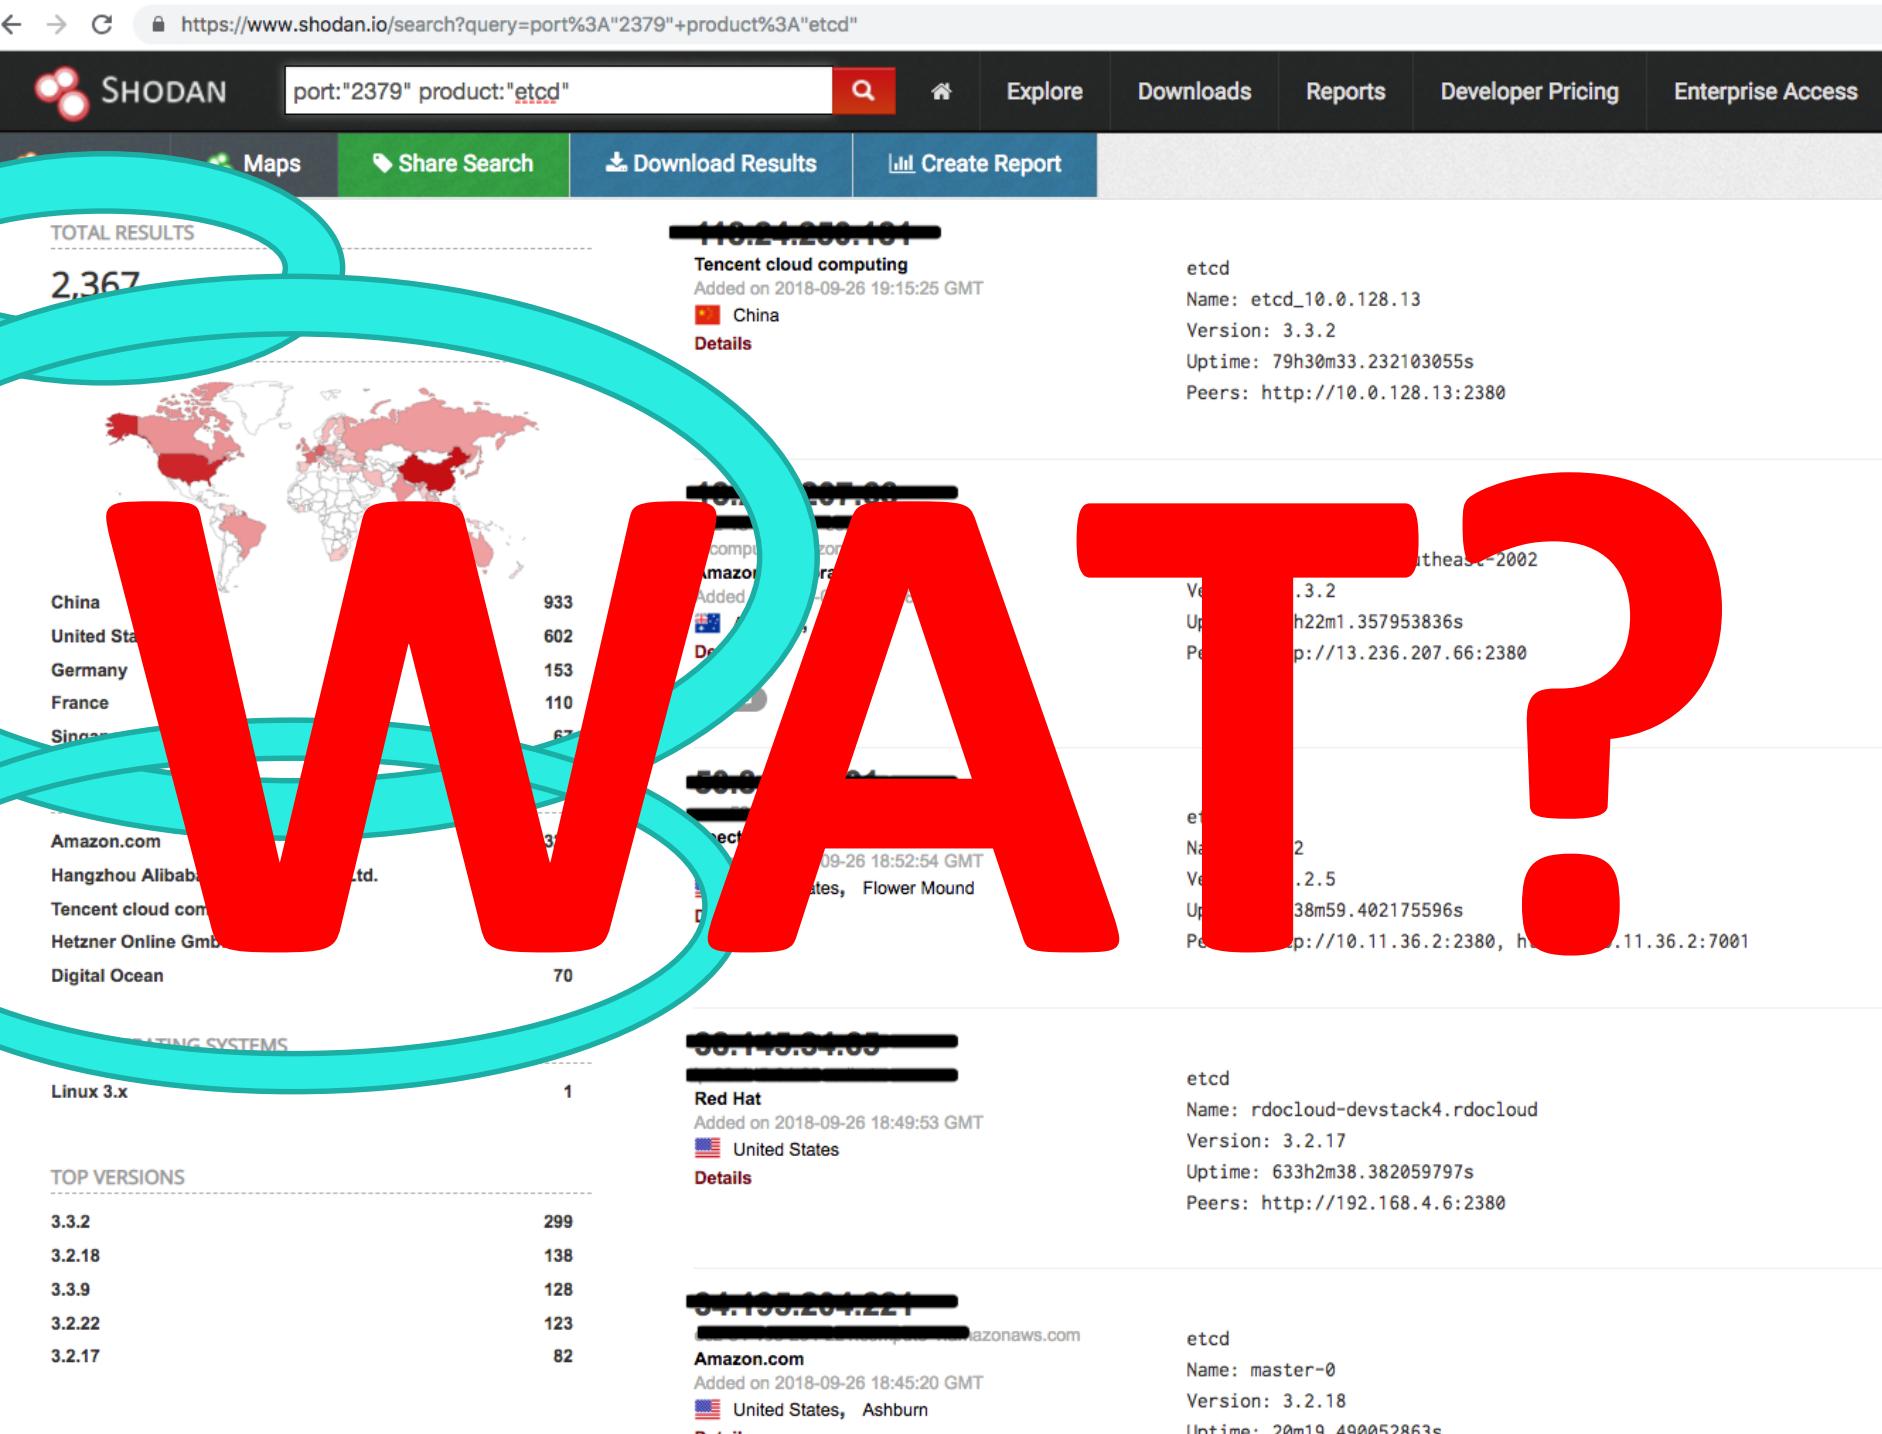
Red Hat
Added on 2018-09-26 18:49:53 GMT
United States
[Details](#)

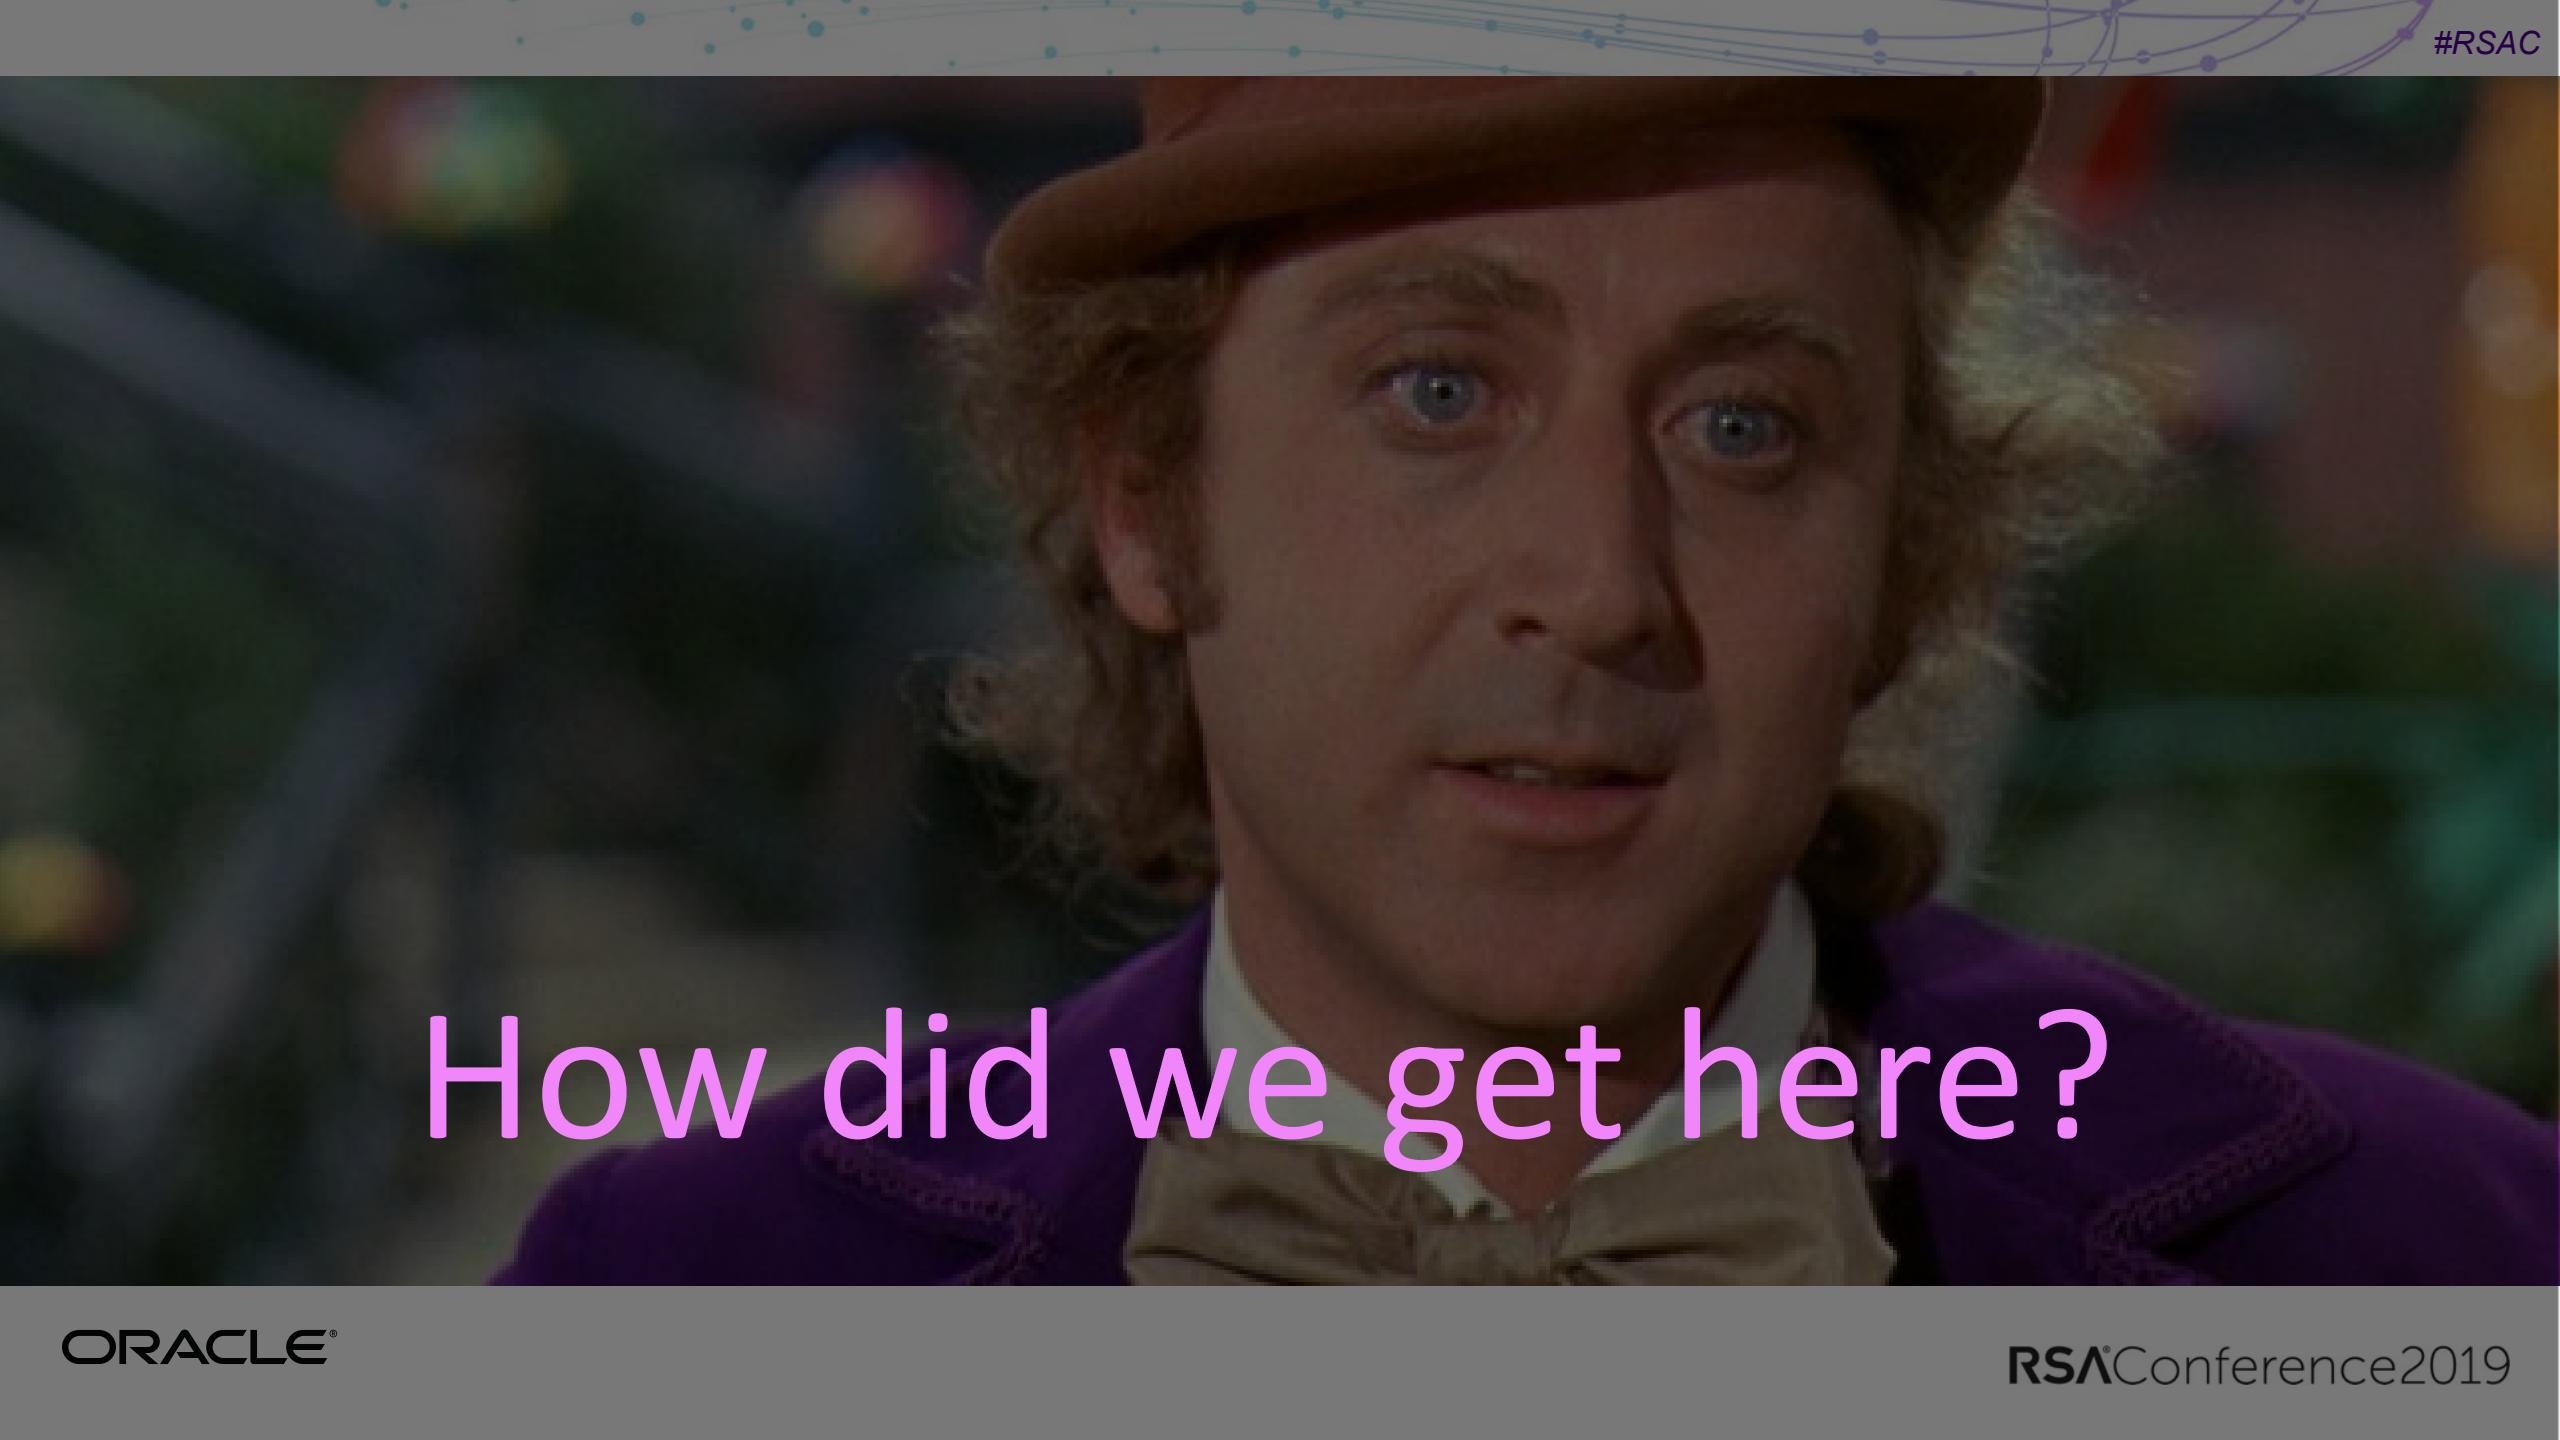
etcd
Name: rdocloud-devstack4.rdocloud
Version: 3.2.17
Uptime: 633h2m38.382059797s
Peers: http://192.168.4.6:2380

54.195.264.221

amazonaws.com
Amazon.com
Added on 2018-09-26 18:45:20 GMT
United States, Ashburn
Digital Ocean

etcd
Name: master-0
Version: 3.2.18
Uptime: 20m19.490052863s





How did we get here?

“Kubernetes is too complicated”



“Kubernetes is too complicated”

“We hope it’ll get easier”



I want to get better!
But where to start?



Let's look at:

- Attack Surface
 - More importantly, how to limit damage
- Security related features in K8s
 - The more you know, the better you build
- Opensource Tooling to help
 - Because we all need help

Attack Surface

Attack Surface

Goal: Reduce the attack surface

- Analysis for:
 - Host(s)
 - Container (Images and running)
 - Kubernetes Cluster

Attack Surface: Host

- These are the machines you're running Kubernetes on.
- Age old principles of Linux still apply:
 - Enable SELinux
 - AppArmor
 - Seccomp
 - Hardened Images
- Goal: Minimize privilege to applications running on the host
- *Good news:* Already a wealth of information on this subject!
 - <http://lmgtfy.com/?q=how+to+reduce+attack+surface+linux>

Attack Surface: Container Images

GOAL: Know your base image when building containers

Attack Surface: Container Images

GOAL: Know your base image when building containers

The screenshot shows a Docker Hub repository page for the public image 'karthequian/ruby'. At the top left is a blue Docker icon. Next to it, the repository name 'karthequian/ruby' is displayed in blue, with 'public' underneath. To the right, there are three tabs: 'STATUS' (highlighted with a teal oval), 'PULLS' (containing the text '500K+ PULLS'), and 'DETAILS'. Below these tabs, the text 'PUBLIC REPOSITORY' is shown in bold. The repository name 'karthequian/ruby' is followed by a star icon. Underneath, the text 'Last pushed: 3 years ago' is visible. At the bottom of the page, there are two tabs: 'Repo Info' (selected) and 'Tags'. The 'Repo Info' section contains fields for 'Short Description' (containing 'A simple sinatra example') and 'Full Description' (containing 'Full description is empty for this repo.'). The 'Tags' section contains a 'Docker Pull Command' field with the value 'docker pull karthequian/ruby' and an 'Owner' field showing a profile picture and the name 'karthequian'.

****BTW, this is just a ruby helloworld app**

Attack Surface: Container Images

GOAL: Know your base image when building containers

The screenshot shows a Docker Hub repository page for the image 'karthequian/ruby'. At the top left is a blue icon of a Docker container. To its right is the repository name 'karthequian/ruby' and the word 'public'. Below this is a teal oval highlighting the repository name 'karthequian/ruby'. To the right of the name is a star icon. Underneath the name, it says 'Last pushed: 3 years ago'. Below the repository name are two tabs: 'Repo Info' (selected) and 'Tags'. The main content area has two columns. The left column contains 'Short Description' (A simple sinatra example) and 'Full Description' (Full description is empty for this repo.). The right column contains 'Docker Pull Command' (docker pull karthequian/ruby) and 'Owner' (karthequian). A teal oval highlights the 'PULLS' count in the top right corner of the repository card, which is labeled '500K+'. Other tabs visible at the top include 'STATUS' and 'DETAILS'.

**BTW, this is just a ruby helloworld app

Attack Surface: Container Images

GOAL: Know your base image when building containers

The screenshot shows a Docker Hub repository page for the image 'karthequian/ruby'. The page includes a public icon, the repository name 'karthequian/ruby' with a star icon, and a status bar showing 0 stars, 500K+ pulls, and details. Below this, there's a 'PUBLIC REPOSITORY' section with tabs for 'Repo Info' and 'Tags'. The 'Repo Info' tab displays a 'Short Description' (A simple sinatra example) and a 'Full Description' (Full description is empty for this repo.). To the right, there's a 'Docker Pull Command' (docker pull karthequian/ruby) and an 'Owner' section (karthequian). Two large cyan ovals highlight the repository name and the full description field.

Attack Surface: Container Images

GOAL: Know your base image when building containers

- When in doubt, stick to an official images!



- Or start from a sane base image (example: alpine linux)

Attack Surface: Container Images

GOAL: Smaller the image, the better

- Less things for an attacker to exploit.
- Quicker to push, quicker to pull.

Attack Surface: Container Images

GOAL: Don't rely on :latest tag

- :latest image yesterday might not be :latest image tomorrow
- Instead, you'd want to know what specific version you're operating with.
- Side benefit: If there is a new vulnerability announced for OS version x.y.z, you know immediately whether you're running that version!

Attack Surface: Container Images

GOAL: Check for vulnerabilities periodically

- Plenty of ways to do this in registries. We'll cover more in the tooling section

Attack Surface: Running Containers

GOAL: Don't run as root

- Containers running as root might be completely unnecessary for the actual application.
- If compromised, attacker can do a lot more things..
- Pod security policies can help (we'll see how later).

Attack Surface: Running Containers

GOAL: Limit host mounts

- Be wary of images that require broad access to paths on the host
- Limit your host mount to a smaller subset of directories
- Reduces blast radius on compromise

Attack Surface: Kubernetes Cluster

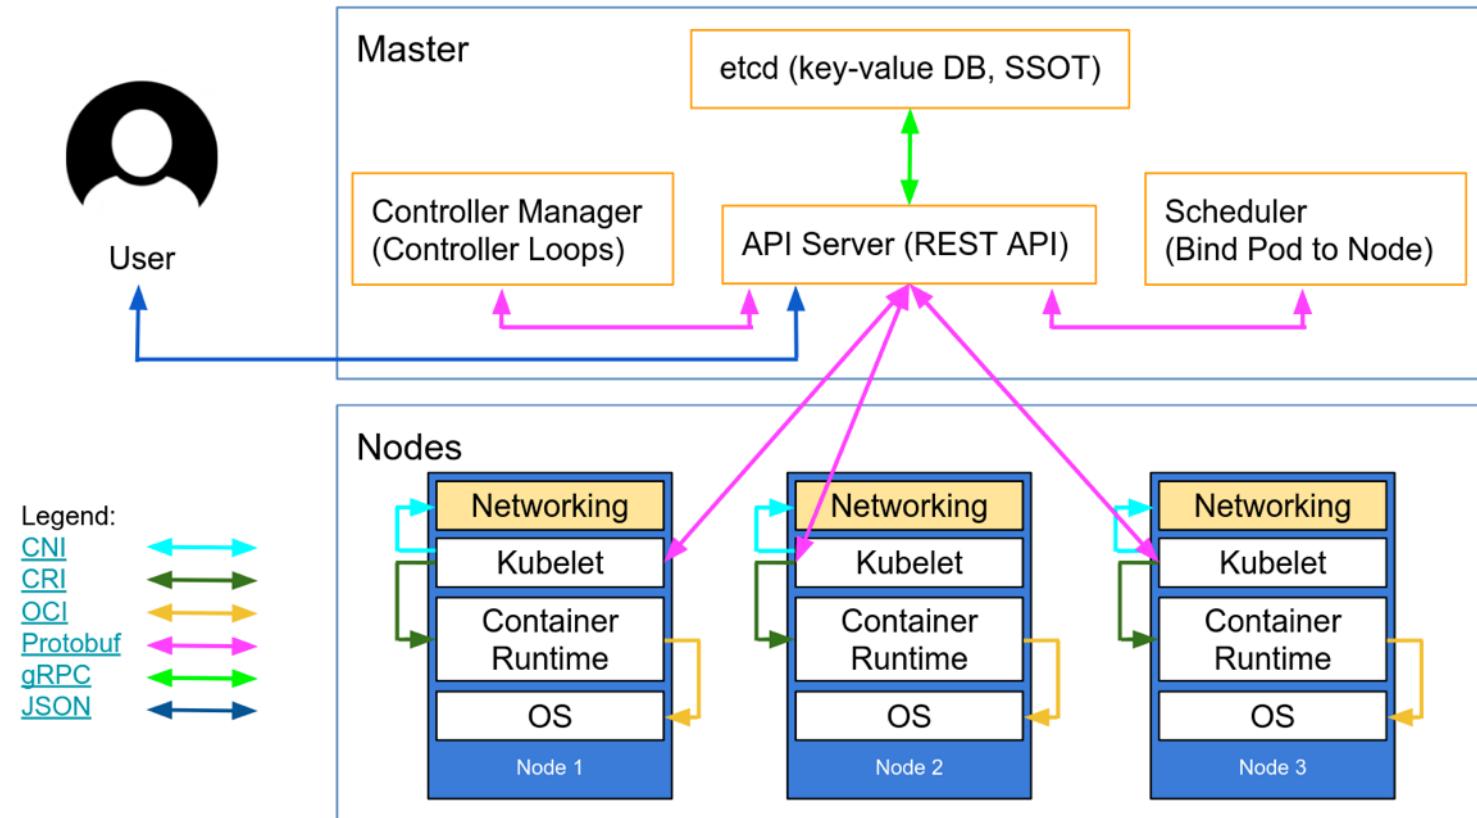
Kubernetes Cluster- TLS

TLS ALL THE THINGS

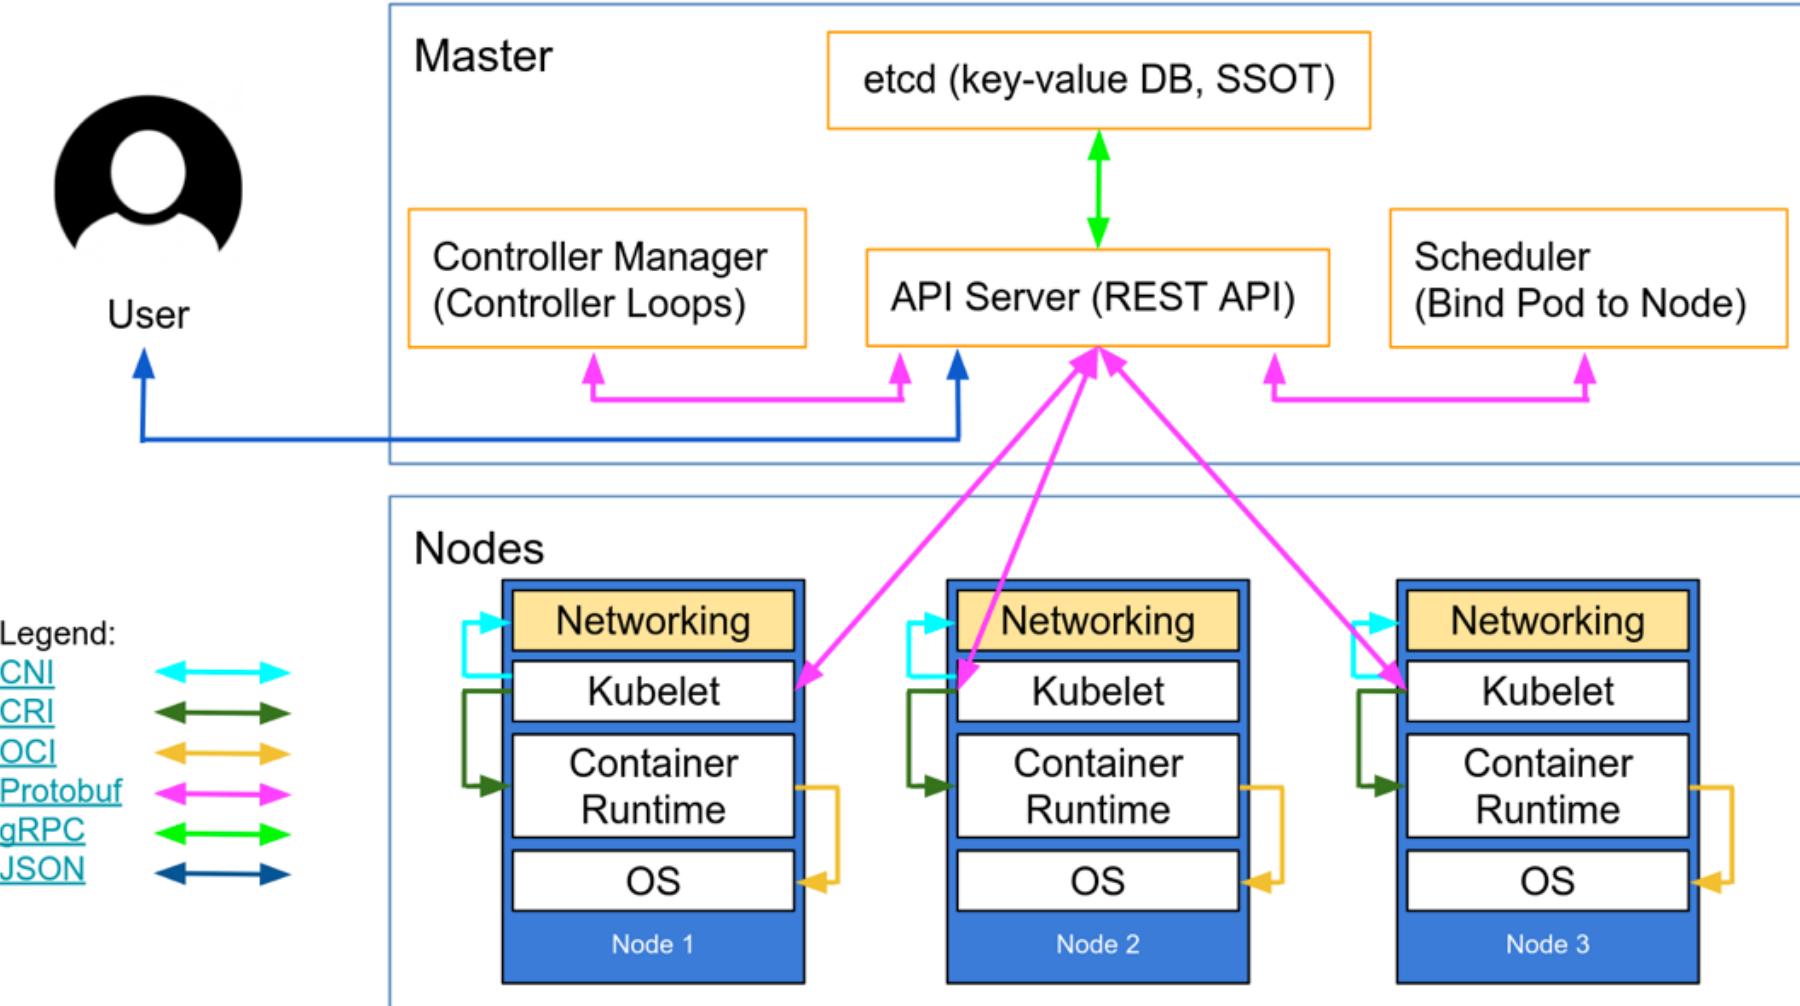


Kubernetes Cluster- TLS

- TLS Checklist:
 1. Nodes and Master
 2. User and Master
 3. Everything etcd
 4. Kubelet to API Server

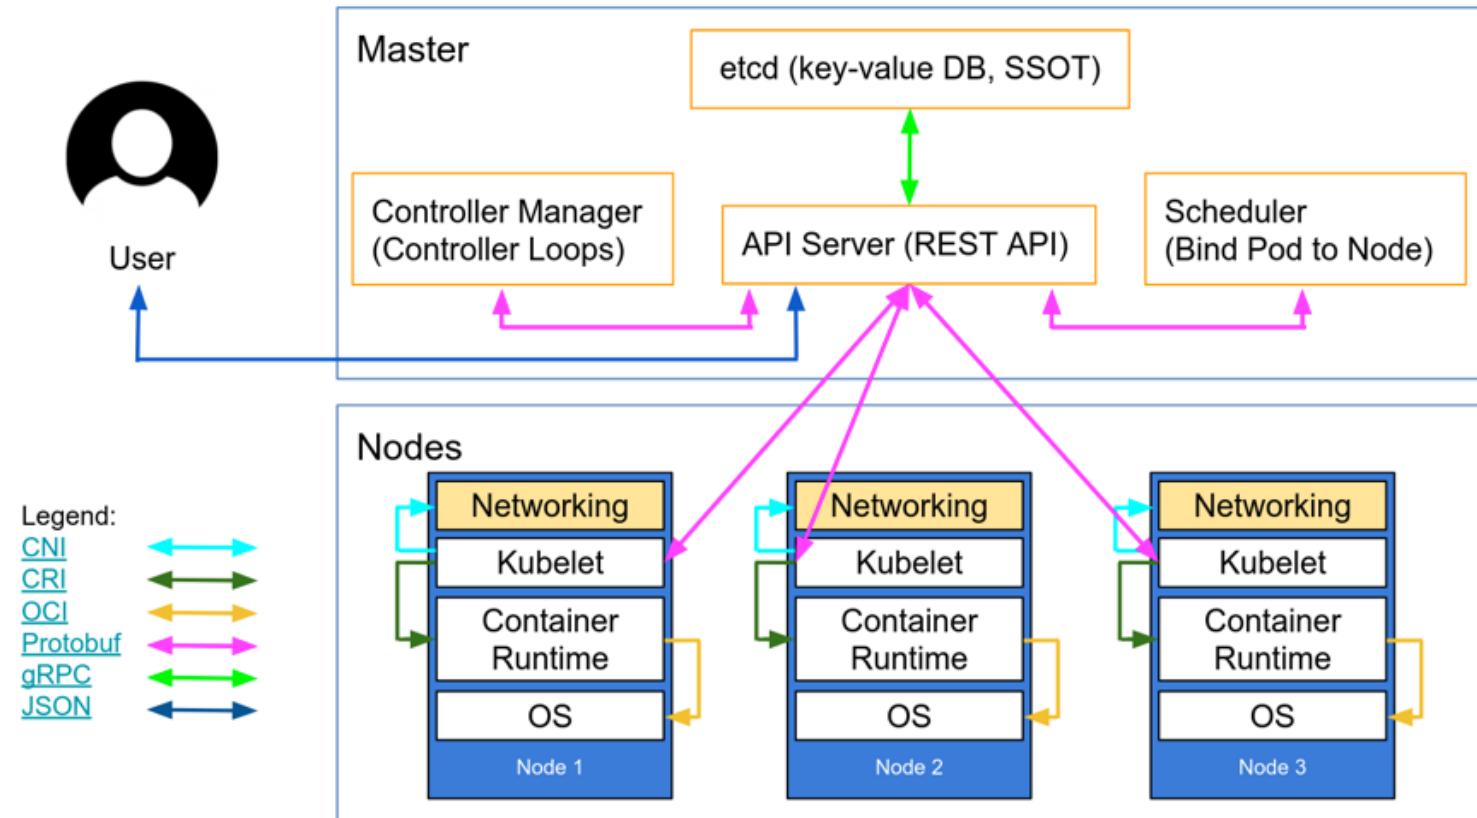


Kubernetes Cluster- TLS



Kubernetes Cluster- TLS

- TLS Checklist:
 1. User and Master
 2. Nodes and Master
 3. Everything etcd
 4. Kubelet to API Server



CVE's Happen...

In a Nutshell, Kubernetes...

... has had [70,892 commits](#) made by [2,241 contributors](#)
representing [1,537,561 lines of code](#)

In a Nutshell, docker...

... has had [257,984 commits](#) made by [11,382 contributors](#)
representing [9,236,254 lines of code](#)

CVE's Happen...

In a Nutshell, Kubernetes...

... has had 70,892 commits made by 2,241 contributors
representing 1,537,561 lines of code

In a Nutshell, docker...

... has had 257,984 commits made by 11,382 contributors
representing 9,236,254 lines of code



CVE-2018-1002105: proxy request handling in kube-apiserver can leave vulnerable TCP connections #71411

DOCKER SECURITY UPDATE: CVE-2019-5736 AND CONTAINER SECURITY BEST PRACTICES

CVE's

GOAL: Have an upgrade strategy

- Because...CVE's are fixed in new minor versions.
- Don't treat K8s as "install once, run all the time".
- Make your K8s install repeatable for different versions.
- ..Or use a Managed Provider.
 - Either automatically patch for you, or tell you what to do.



We're a little
better off now.

But what else to do?

K8s Features

How can the platform help
me make secure choices?



K8s Features

- Authentication
- Authorization
- Audit Logging
- Network Policies
- Pod security policies
- Kubernetes Secrets



Authentication and Authorization

- Do you know how you are authenticating with Kubernetes?
- Many ways to Authenticate
 - Client Certs
 - Static token file
 - Service Account tokens
 - OpenID
 - Webhook Mode
 - And more (<https://kubernetes.io/docs/reference/access-authn-authz/authentication/>)



Goal: Pick a strategy that fits
your use case

Whatever you do,
DO NOT YOLO!

A close-up photograph of a fluffy, light-colored cat with large, wide-open eyes, looking directly at the camera. The cat is positioned behind a dark, textured surface, possibly a book or a piece of furniture, which is visible in the background. The lighting is soft, highlighting the cat's fur and eyes.

If you DO NOT YOLO...

You can pick an authz strategy..

Authentication and Authorization

Authorization Modules

- **Node** - A special-purpose authorizer that grants permissions to kubelets based on the pods they are scheduled to run. To learn more about using the Node authorization mode, see [Node Authorization](#).
- **ABAC** - Attribute-based access control (ABAC) defines an access control paradigm whereby access rights are granted to users through the use of policies which combine attributes together. The policies can use any type of attributes (user attributes, resource attributes, object, environment attributes, etc). To learn more about using the ABAC mode, see [ABAC Mode](#).
- **RBAC** - Role-based access control (RBAC) is a method of regulating access to computer or network resources based on the roles of individual users within an enterprise. In this context, access is the ability of an individual user to perform a specific task, such as view, create, or modify a file. To learn more about using the RBAC mode, see [RBAC Mode](#)
 - When specified RBAC (Role-Based Access Control) uses the `rbac.authorization.k8s.io` API group to drive authorization decisions, allowing admins to dynamically configure permission policies through the Kubernetes API.
 - To enable RBAC, start the apiserver with `--authorization-mode=RBAC`.
- **Webhook** - A WebHook is an HTTP callback: an HTTP POST that occurs when something happens; a simple event-notification via HTTP POST. A web application implementing WebHooks will POST a message to a URL when certain things happen. To learn more about using the Webhook mode, see [Webhook Mode](#).

<https://kubernetes.io/docs/reference/access-authn-authz/authorization/>

Authentication and Authorization

- Pro tip: Nobody uses ABAC anymore. Don't be that guy....
- RBAC is the defacto standard
 - Based on roles and role bindings
 - Good set of defaults: <https://github.com/uruddaraju/kubernetes-rbac-policies>
- Can use multiple authorizers together, but can get confusing.
 - 1st authorizer to authorize passes authz

Kubernetes Cluster- Audit Logs

- Wat?
- “Kubernetes auditing provides a security-relevant chronological set of records documenting the sequence of activities that have affected system by individual users, administrators or other components of the system.”
- Answers: What/when/who/where information on security events.
- **Your job:** Periodically watch Kubernetes Audit logs
- <https://kubernetes.io/docs/tasks/debug-application-cluster/audit/>



See, you know how to take the reservation, you just don't know how to hold the reservation and that's really the most important part of the reservation, the holding.
Anybody can just take them.

— *Jerry Seinfeld* —

AZ QUOTES

Kubernetes Cluster- Network Policies

- Consider adding a network policy to the cluster...
- Default Policy: All pods can talk to all other pods.
- Consider limiting this with a Network Policy
- <https://kubernetes.io/docs/concepts/services-networking/network-policies/>

Kubernetes Cluster- Pod Security Policies

- Consider adding Pod Security policies
- PodSecurityPolicy: A Defined set of conditions a pod must run with.
- Think of this as authorization for pods.

Kubernetes Cluster: Pod Security Policies

Capability for an admin to control specific actions

Control Aspect	Field Names
Running of privileged containers	privileged
Usage of host namespaces	hostPID , hostIPC
Usage of host networking and ports	hostNetwork , hostPorts
Usage of volume types	volumes
Usage of the host filesystem	allowedHostPaths
White list of Flexvolume drivers	allowedFlexVolumes
Allocating an FSGroup that owns the pod's volumes	fsGroup
Requiring the use of a read only root file system	readOnlyRootFilesystem
The user and group IDs of the container	runAsUser , supplementalGroups
Restricting escalation to root privileges	allowPrivilegeEscalation , defaultAllowPrivilegeEscalation
Linux capabilities	defaultAddCapabilities , requiredDropCapabilities , allowedCapabilities
The SELinux context of the container	selinux
The AppArmor profile used by containers	annotations
The seccomp profile used by containers	annotations
The sysctl profile used by containers	annotations

<https://kubernetes.io/docs/concepts/policy/pod-security-policy/#what-is-a-pod-security-policy>

Kubernetes Secrets

- **GOAL: Use Kubernetes secrets to store sensitive data instead of config maps.**
- Also look at: secrets encryption provider.
 - Controls how etcd encrypts API data
 - **--experimental-encryption-provider-config**
- <https://kubernetes.io/docs/tasks/administer-cluster/encrypt-data/>

Opensource Tooling



Keep tabs on the CNCF Security landscape

CNCF Cloud Native Interactive Landscape

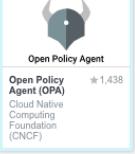
CNCF's Cloud Native Trail Map ([png](#), [pdf](#)) provides a good introduction. The cloud native landscape ([png](#), [pdf](#)) and serverless landscape ([png](#), [pdf](#)) are dynamically generated below. Please [open](#) a pull request to correct any issues. Greyed logos are not open source. Last Updated: 2018-09-26 04:01:06Z
You are viewing 20 cards with a total of 16,396 stars, market cap of \$235B and funding of \$455M.

[CARD MODE](#) [LANDSCAPE](#) [SERVERLESS](#)

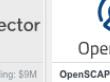
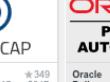
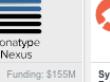
Incubating CNCF Projects (2)

	Notary Cloud Native Computing Foundation (CNCF) ★ 1,527		The Update Framework (TUF) Cloud Native Computing Foundation (CNCF) ★ 785
---	---	---	---

Sandbox CNCF Projects (1)

	Open Policy Agent Cloud Native Computing Foundation (CNCF) ★ 1,438
---	--

CNCF Member Products/Projects (11)

	Aqua Aqua Security Funding: \$38M		clair Red Hat MCap: \$24B ★ 4,173		Datica Datica Funding: \$12.8M		kube-bench Aqua Security Funding: \$88M		kube-hunter Aqua Security Funding: \$38M ★ 574		NeuVector NeuVector Funding: \$9M		OpenSCAP Red Hat MCap: \$24B ★ 349		ORACLE POLICY AUTOMATION Oracle MCap: \$196B		Sonatype Nexus Sonatype Funding: \$155M		Sysdig Falco Sysdig ★ 938 Funding: \$122M		Twistlock Twistlock Funding: \$63.1M
---	---	---	--	---	--------------------------------------	---	---	--	---	---	---	---	---	---	--	---	---	---	--	---	--

Non-CNCF Member Products/Projects (6)

	anchore Anchore Funding: \$273		BLACKDUCK Black Duck Synopsys MCap: \$14.7B		Grafeas Grafeas ★ 631		ORY / Hydra ORY Hydra ★ 4,825		StackRox StackRox Funding: \$39M		WhiteSource WhiteSource Funding: \$11M
---	--------------------------------------	---	--	---	-----------------------------	---	--	--	--	---	--

Crunchbase data is used under license from Crunchbase to CNCF. For more information, please see the [license](#) info.

<https://landscape.cncf.io/landscape=security-compliance>

CNCF Projects



- “The Update Framework”
- Is a framework or a methodology.
- Used for secure software updates.
- Based on ideas surrounding trust and integrity.

- Is a project.
- Based on TUF.
- A solution to secure software updates and distribution.
- Used in Docker Trusted Registry.

Clair

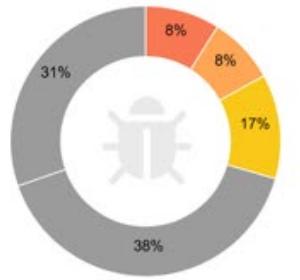


clair

- Open source project for the static analysis of vulnerabilities in containers.
- Find vulnerable images in your repo.
- Built into quay.io, but you can add to your own repo.
- <https://github.com/coreos/clair>

← example/repository

56c0fa71e47b



Quay Security Scanner has detected **13** vulnerabilities.

Patches are available for **4** vulnerabilities.

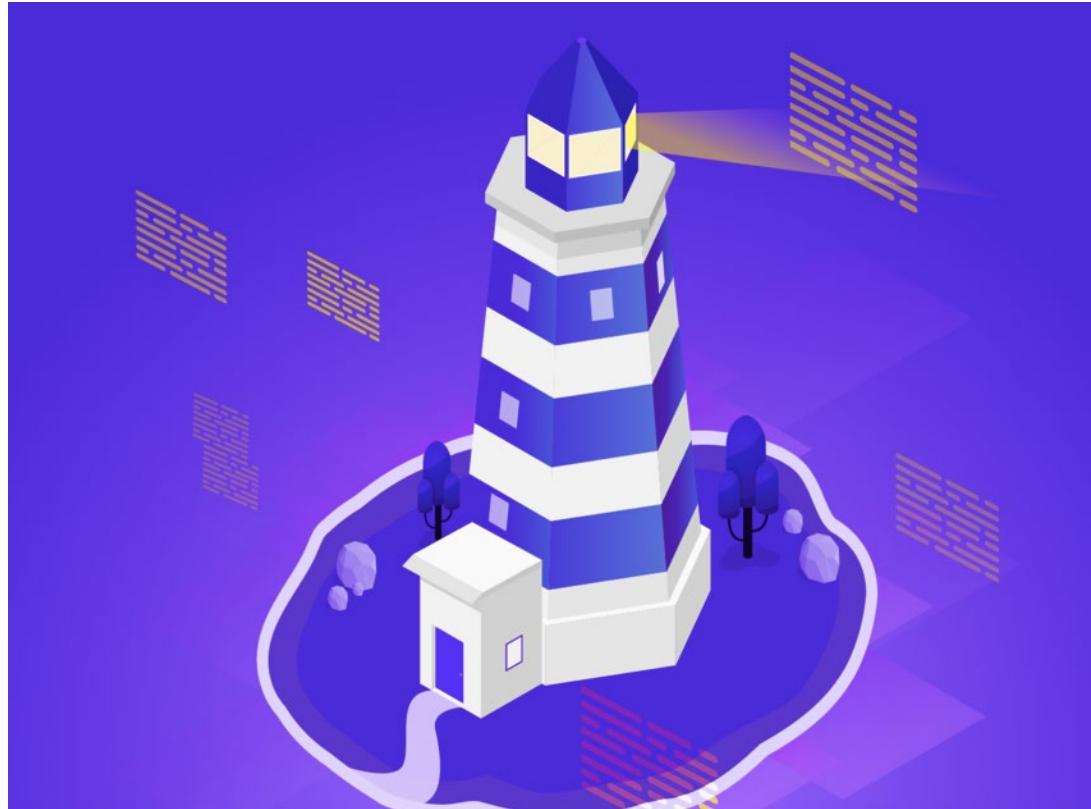
- ⚠ 1 High-level vulnerabilities.
- ⚠ 1 Medium-level vulnerabilities.
- ⚠ 2 Low-level vulnerabilities.
- ⚠ 5 Negligible-level vulnerabilities.
- ⚠ 4 Unknown-level vulnerabilities.

Image Vulnerabilities

 Filter Vulnerabilities... Only show fixable

CVE	Severity	Package	Current Version	Fixed In Version	Introduced In Image	Action
▶ CVE-2013-7445	⚠ High	linux	3.16.7-ckt20-1+deb8u3	(None)		apt-get up.
▶ CVE-2015-5276	⚠ Medium	gcc-4.9	4.9.2-10	(None)		file:b5391.
▶ CVE-2016-2856	⚠ Low	glibc	2.19-18+deb8u3	(None)		file:b5391.
▶ CVE-2016-0823	⚠ Low	linux	3.16.7-ckt20-1+deb8u3	(None)		apt-get up.
▶ CVE-2005-3660	⚠ Negligible *	linux	3.16.7-ckt20-1+deb8u3	(None)		apt-get up.
▶ CVE-2015-4003	⚠ Negligible *	linux	3.16.7-ckt20-1+deb8u3	(None)		apt-get up.
▶ CVE-2008-4108	⚠ Negligible *	python-defaults	2.7.9-1	(None)		apt-get up.
▶ CVE-2015-8830	⚠ Negligible	linux	3.16.7-ckt20-1+deb8u3	3.16.7-ckt20-1+deb8u4		apt-get up.
▶ CVE-2013-4392	⚠ Negligible *	systemd	215-17+deb8u3	(None)		file:b5391.
▶ CVE-2015-7515	⚠ Unknown	linux	3.16.7-ckt20-1+deb8u3	(None)		apt-get up.
▶ CVE-2015-8816	⚠ Unknown	linux	3.16.7-ckt20-1+deb8u3	3.16.7-ckt20-1+deb8u4		apt-get up.
▶ CVE-2016-2547	⚠ Unknown	linux	3.16.7-ckt20-1+deb8u3	3.16.7-ckt20-1+deb8u4		apt-get up.
▶ CVE-2016-2545	⚠ Unknown	linux	3.16.7-ckt20-1+deb8u3	3.16.7-ckt20-1+deb8u4		apt-get up.

Harbor



- Newer! CNCF Project
- Registry product
- Supports vulnerability scanning, image signing and identity control
- Scope is larger than clair

Harbor

Harbor

 Search Harbor...

English

About

< Projects < Repositories < thomas/postgres

thomas/postgres:alpine

Author	anonymity
Architecture	amd64
OS	linux
Docker Version	17.06.2-ce
Scan Completed	Nov 2, 2018



SCAN

Vulnerability	Severity	Package	Current version	Fixed in version
> CVE-2018-9251	Low	libxml2	2.9.8-r0	2.9.8-r1
> CVE-2018-14404	Medium	libxml2	2.9.8-r0	2.9.8-r1
> CVE-2018-14567	Medium	libxml2	2.9.8-r0	2.9.8-r1

Kube-bench

- Checks whether a Kubernetes cluster is deployed according to security best practices.
- Run this after creating your K8s cluster.
- <https://github.com/aquasecurity/kube-bench>
- Defined by the CIS Benchmarks Docs:
<https://www.cisecurity.org/cis-benchmarks/>
- Run it against your Kubernetes Master, or Kubernetes node.



kube-bench

Kube-bench example

```
› ~$ kubectl logs kube-bench-node
[INFO] 2 Worker Node Security Configuration
[INFO] 2.1 Kubelet
[FAIL] 2.1.1 Ensure that the --allow-privileged argument is set to false (Scored)
[PASS] 2.1.2 Ensure that the --anonymous-auth argument is set to false (Scored)
[PASS] 2.1.3 Ensure that the --authorization-mode argument is not set to AlwaysAllow (Scored)
[PASS] 2.1.4 Ensure that the --client-ca-file argument is set as appropriate (Scored)
[PASS] 2.1.5 Ensure that the --read-only-port argument is set to 0 (Scored)
[FAIL] 2.1.6 Ensure that the --streaming-connection-idle-timeout argument is not set to 0 (Scored)
[FAIL] 2.1.7 Ensure that the --protect-kernel-defaults argument is set to true (Scored)
[FAIL] 2.1.8 Ensure that the --make-iptables-util-chains argument is set to true (Scored)
[FAIL] 2.1.9 Ensure that the --keep-terminated-pod-volumes argument is set to false (Scored)
[FAIL] 2.1.10 Ensure that the --hostname-override argument is not set (Scored)
[FAIL] 2.1.11 Ensure that the --event-qps argument is set to 0 (Scored)
[PASS] 2.1.12 Ensure that the --tls-cert-file and --tls-private-key-file arguments are set as appropriate (Scored)
[PASS] 2.1.13 Ensure that the --cadvisor-port argument is set to 0 (Scored)
[FAIL] 2.1.14 Ensure that the RotateKubeletClientCertificate argument is set to true
[FAIL] 2.1.15 Ensure that the RotateKubeletServerCertificate argument is set to true
[INFO] 2.2 Configuration Files
[FAIL] 2.2.1 Ensure that the kubelet.conf file permissions are set to 644 or more restrictive (Scored)
[FAIL] 2.2.2 Ensure that the kubelet.conf file ownership is set to root:root (Scored)
[FAIL] 2.2.3 Ensure that the kubelet service file permissions are set to 644 or more restrictive (Scored)
[FAIL] 2.2.4 Ensure that the kubelet service file ownership is set to root:root (Scored)
[FAIL] 2.2.5 Ensure that the proxy kubeconfig file permissions are set to 644 or more restrictive (Scored)
[FAIL] 2.2.6 Ensure that the proxy kubeconfig file ownership is set to root:root (Scored)
[WARN] 2.2.7 Ensure that the certificate authorities file permissions are set to 644 or more restrictive (Scored)
[WARN] 2.2.8 Ensure that the client certificate authorities file ownership is set to root:root
```

Kubesec

- Helps you quantify risk for Kubernetes resources.
- Run against your K8s applications
(deployments/pods/daemonsets etc)
- <https://kubesec.io/> from controlplane
- Can be used standalone, or as a kubectl plugin
(<https://github.com/stefanprodan/kubectl-kubesec>)

Kubesec example

```
▶ ~$ kubectl -n kube-system plugin scan deployment/kubernetes-dashboard  
scanning deployment kubernetes-dashboard  
deployment/kubernetes-dashboard kubesec.io score 3
```

Advise

1. containers[] .securityContext .runAsNonRoot == true

Force the running image to run as a non-root user to ensure least privilege

2. containers[] .securityContext .capabilities .drop

Reducing kernel capabilities available to a container limits its attack surface

3. containers[] .securityContext .readOnlyRootFilesystem == true

An immutable root filesystem can prevent malicious binaries being added to PATH and increase attack cost

4. containers[] .securityContext .runAsUser > 10000

Run as a high-UID user to avoid conflicts with the host's user table

5. containers[] .securityContext .capabilities .drop | index("ALL")

Drop all capabilities and add only those required to reduce syscall attack surface

```
▶ ~$ █
```



Kubeaudit

- Opensourced from Shopify.
- Auditing your applications in your K8s cluster.
- <https://github.com/Shopify/kubeaudit>
- Little more targeted than Kubesec.

kubeaudit is a program that will help you audit your Kubernetes clusters. Specify -l to run kubeaudit using ~/.kube/config otherwise it will attempt to create an in-cluster client.

#patcheswelcome

Usage:

```
kubeaudit [command]
```

Available Commands:

allowpe	Audit containers that allow privilege escalation
caps	Audit container for capabilities
help	Help about any command
image	Audit container images
nonroot	Audit containers running as root
np	Audit namespace network policies
priv	Audit containers running as root
rootfs	Audit containers with read only root filesystems
sat	Audit automountServiceAccountToken = true pods against an empty (default) service account
version	Print the version number of kubeaudit

Flags:

-a, --allPods	Audit againsts pods in all the phases (default Running Phase)
-h, --help	help for kubeaudit
-j, --json	Enable json logging
-c, --kubeconfig string	config file (default is \$HOME/.kube/config)
-l, --local	Local mode, uses ~/.kube/config as configuration
-f, --manifest string	yaml configuration to audit
-v, --verbose string	Set the debug level (default "INFO")

Kubeaudit example

```
~$ ./Users/karthik/Downloads/kubeaudit_0.2.0_darwin_amd64/kubeaudit allowpe -c ./Users/karthik/.kube/config
ERRO[0003] SecurityContext not set, please set it! KubeType=deployment Name=dotnetworld Namespace=default
ERRO[0003] SecurityContext not set, please set it! KubeType=deployment Name=javademo Namespace=default
ERRO[0003] SecurityContext not set, please set it! KubeType=deployment Name=prom-demo-prometheus-alertmanager Namespace=default
ERRO[0003] SecurityContext not set, please set it! KubeType=deployment Name=prom-demo-prometheus-kube-state-metrics Namespace=default
ERRO[0003] SecurityContext not set, please set it! KubeType=deployment Name=prom-demo-prometheus-pushgateway Namespace=default
ERRO[0003] SecurityContext not set, please set it! KubeType=deployment Name=prom-demo-prometheus-server Namespace=default
ERRO[0003] SecurityContext not set, please set it! KubeType=deployment Name=wishlist-deployment Namespace=default
ERRO[0003] SecurityContext not set, please set it! KubeType=deployment Name=contour Namespace=heptio-contour
ERRO[0003] SecurityContext not set, please set it! KubeType=deployment Name=kube-dns Namespace=kube-system
ERRO[0003] SecurityContext not set, please set it! KubeType=deployment Name=kube-dns-autoscaler Namespace=kube-system
ERRO[0003] SecurityContext not set, please set it! KubeType=deployment Name=kubernetes-dashboard Namespace=kube-system
ERRO[0003] SecurityContext not set, please set it! KubeType=deployment Name=oci-volume-provisioner Namespace=kube-system
ERRO[0003] SecurityContext not set, please set it! KubeType=deployment Name=tiller-deploy Namespace=kube-system
ERRO[0003] SecurityContext not set, please set it! KubeType=daemonSet Name=prom-demo-prometheus-node-exporter Namespace=default
ERRO[0003] AllowPrivilegeEscalation not set which allows privilege escalation, please set to false KubeType=daemonSet Name=kube-flannel-ds Namespace=kube-system
ERRO[0003] AllowPrivilegeEscalation not set which allows privilege escalation, please set to false KubeType=daemonSet Name=kube-proxy Namespace=kube-system
```

Couple more resources to look at:

- 11 ways not to get hacked:

<https://kubernetes.io/blog/2018/07/18/11-ways-not-to-get-hacked>

- K8s security (from Image Hygiene to Network Policy):

<https://speakerdeck.com/mhausenblas/kubernetes-security-from-image-hygiene-to-network-policies>

Apply It!



Apply It!

- Day 1:
- Know what version of Docker and Kubernetes you use.
- Understand if your control and data plane nodes are hardened.
- Understand how your Docker containers are built.
- Find out how you authenticate and authorize for your clusters.



KNOWING
IS HALF THE BATTLE

Apply It!

- Week 1:
- Build an Automation Pipeline:
 - To build Docker images on code pushes
 - Versioning strategy for code
 - To build your Kubernetes clusters

Apply It!

- **1st Month**
- Sanitize your code:
 - Know your base images
 - Implement versioning for your containers
 - Invest in a registry (or tooling) that does vulnerability scanning
- Kubernetes:
 - Have an upgrade strategy in place
 - Analyze secrets/sensitive cluster data
 - Turn on audit logging

Apply It!

- 3 Months:
- Continuously Monitor
 - Tooling like Kubesec/Kube-audit
- Plan how to address vulnerabilities/CVE's
- K8s:
 - Strategy for Pod Security Policies
 - Strategy for Network Policies
 - Run scans (like kube-bench) on cluster creation

Apply It!

- 6 Months:
- Re-ask day 1 questions.
- Review strategies- is it working? What needs tweaking?
- Review tooling- are there new tools that help?
- Review CVE's

KEEP CALM
AND
KUBE ON

@iteration1

